

FROM FEEDBACK TO MEMORY TO BEHAVIOR: HOW PERSISTENT DIFFERENCES ARE WRITTEN AND WHERE NATIVE TRANSPORT BREAKS

Anonymous authors

Paper under double-blind review

ABSTRACT

Self-improving agents often turn terminal feedback into persistent natural-language memory, creating a two-part scientific problem: *how does feedback write different persistent state from the same experience, and why does that state difference not necessarily become a behavioral difference later?* We study both parts with frozen success-versus-failure writer branches on otherwise identical trajectories, while treating the released branch switch as a bundled writer-protocol intervention rather than a pure reward-bit manipulation. First, the write channel is robust: the intervention produces different persistent memories for all 20 Shopping pairs and all 4 Reddit pairs, and a stronger same-mode wording control still leaves a significant branch-mode excess. Second, the downstream story is selective. When branch-specific memory is mechanically supplied, it has measurable terminal leverage ($|\Delta| = 0.15625$, $p = 0.00074$). Under native reuse, however, Shopping retrieves the branch-conditioned source-memory item in 125/172 opportunities while the first-action branch contrast is not supported ($TV = 0.06944$, $p = 0.5801$, 0/36 modal changes) and terminal transport is sparse ($|\Delta| = 0.02083$, $p = 0.4289$, 34/36 zero). Reddit again shows reliable write divergence but sparse, sign-heterogeneous native terminal differences. We therefore follow the same persistent-state contrast through write, native exposure, first-action uptake, and outcome, with forced fixed-evidence exposure as a separate capacity control. The resulting stage signature reconciles an apparent contradiction: feedback-conditioned writing can reliably create different durable memory, and that difference can have behavioral leverage when supplied, yet the deployed native reuse chain need not transport it into stable decision-level change. On the frozen Shopping tests, action uptake is the first unsupported measured native stage after write and item-level exposure have been established. Because stage units and test sensitivities are not matched, this is an evidence-localization result rather than a causal mediation estimate. The paper thus connects two questions usually studied separately: how evaluative feedback becomes persistent agent state, and where that state loses behavioral authority during reuse.

1 INTRODUCTION

A persistent-memory agent can diverge from itself twice: first when feedback changes what it stores from an otherwise identical experience, and again when only some of that stored difference survives into later behavior. Consider two copies of the same interaction trajectory. If one copy is routed through the released success-reflection branch and the other through the failure-reflection branch, the writer can store different reusable advice even though the underlying experience is unchanged. The first scientific question is therefore *how strongly does feedback-conditioned writing change persistent state?* The second is *once that difference exists, where does the deployed reuse chain preserve or lose it before behavior changes?*

These two questions belong in the same paper because persistent self-improvement contains several causal surfaces that are often collapsed into one “memory helped” statement. A terminal signal may first alter what is **written**; the resulting branch-conditioned source-memory item may or may not be **exposed** by retrieval on a later task; the policy may or may not **use** the treatment-specific information

within that retrieved item; and only then can behavior change at the terminal **outcome**. A positive memory-similarity result identifies the first surface, while a retrieval hit only identifies availability. Neither is evidence by itself that the agent’s decision rule changed. Conversely, an endpoint null does not reveal whether the update failed to write, failed to retrieve, or was filtered after retrieval. We call this *evaluation coarsening*: observing only a projection of the stage chain aliases distinct internal explanations into the same coarse success or failure label.

We study this problem through a stage-resolved intervention rather than a new memory algorithm. For a frozen trajectory τ , we switch the released success/failure reflection branch and construct paired persistent memories while holding the source experience fixed. This is a bundled writer-protocol intervention: the branch changes the reflection instruction together with its outcome/reward semantics, so it is not an atom-level isolation of a reward bit. We then instrument four native stages,

write $\rightarrow$ retrieval exposure $\rightarrow$ first-action uptake $\rightarrow$ terminal outcome,

and add a separate forced fixed-evidence control that bypasses native retrieval. The forced arm is deliberately *not* another stage in the chain: it asks whether branch-specific memory content has downstream leverage when supplied, whereas the native path asks whether the deployed memory system actually transports that difference into behavior.

This design distinguishes several explanations that an endpoint-only experiment cannot separate. If the writer does not create a persistent intervention, the story stops at write time. If the downstream policy is globally insensitive to memory content, even forced exposure should have little leverage. If the only problem is retrieval absence, native exposure should collapse. If retrieval itself is a valid proxy for policy use, observed exposure should co-occur with a supported action-level branch contrast. Finally, if reward-conditioned memory has a universal directional effect, terminal differences should not be dominated by zeros and sign reversals across tasks. Because these stages have different units, sample units, tests, and experimental surfaces, we do not convert them into a single “transport efficiency.” Instead, we retain a *stage signature* and introduce an ordinal stage-evidence ladder: preserve each stage’s own inferential/descriptive semantics, then localize the first *measured* native stage whose branch contrast is no longer supported after the preceding prerequisites have been supported or directly observed. The stage tests are not power-matched, so this operator localizes an evidence boundary rather than the latent point where a causal effect must physically decay. The point is diagnostic disambiguation rather than a finer-grained score.

The frozen evidence yields a sharp pattern. Reward-conditioned writing changes durable memory in all 20 Shopping pairs and all 4 Reddit pairs. A same-mode wording control shows that the reward/reflection-mode contrast exceeds a stronger within-mode wording perturbation. Forced fixed-evidence exposure produces terminal $|\Delta| = 0.15625$ ($p = 0.00074$), demonstrating downstream leverage under supplied memory. Yet native Shopping retrieval remains common—125 hits in 172 held-out opportunities—while first-action TV is only 0.06944 ($p = 0.5801$) with 0/36 modal action changes, and native terminal $|\Delta|$ is 0.02083 ($p = 0.4289$) with 34/36 zero cells. Reddit again shows reliable write divergence but sparse terminal transport: 6/8 cells are zero and the two nonzero effects have opposite signs.

Taken together, these measurements weaken three simple explanations at once: the native null is not caused by failure to create persistent state; it is not well described as global downstream insensitivity to supplied memory content; and it is not explained by retrieval absence alone. More importantly, they expose why coarse evaluations disagree. Write-only evaluation would call the update strong; retrieval-only evaluation would call reuse active; native endpoint-only evaluation would call the effect negligible; forced-only evaluation would call the memory behaviorally consequential. The full stage signature resolves this apparent contradiction: on the frozen Shopping tests, first-action uptake is the first unsupported measured native stage after durable write and source-item exposure have been established. We intentionally stop short of a causal mediation claim because the frozen pool lacks per-atom decoding-seed binding and same-condition noise-floor replication, and the forced and native surfaces are not identical interventions.

Our contributions are:

1. **A controlled feedback-to-state result.** Holding the source trajectory fixed, we switch the released success/failure reflection branch as a bundled writer-protocol intervention and show robust persistent-state divergence across Shopping and Reddit, including a stronger same-mode wording control that does not explain the write effect away.

2. **A stage-resolved account of what happens after writing.** We follow that same persistent-state contrast through native source-item exposure, first-action uptake, and terminal outcome rather than treating state divergence as behavioral authority.
3. **A capacity-versus-transport control.** Forced fixed-evidence exposure establishes that branch-specific memory can have downstream leverage when supplied, while native reuse tests whether the deployed retrieval-and-policy path actually realizes that leverage.
4. **A diagnostic stage signature and cross-domain boundary.** We show that write-only, retrieval-only, native-endpoint-only, and forced-only views support different coarse summaries, and localize the first unsupported measured native Shopping stage without inventing a cross-stage attenuation ratio or causal mediation coefficient. Reddit reproduces write-stage divergence but not a stable signed native terminal effect.

The resulting claim is deliberately narrower than a new memory method: **persistent-state divergence is not behavioral authority, and retrieval availability is not policy use**. This boundary tells us what a self-improving agent evaluation must measure before crediting a persistent update with downstream control.

2 STAGE-RESOLVED MEMORY TRANSPORT

2.1 CONTROLLED WRITE INTERVENTION

Let τ denote a frozen interaction trajectory before persistent memory construction. We create two writer conditions on the same τ ,

$$r \in \{S, F\}, \quad m^{(r)} = g(\tau, r),$$

where g is the released memory writer and S/F selects its success- or failure-reflection mode. The operational write contrast is

$$W(\tau) = d_{\text{mem}}(m^{(S)}, m^{(F)}).$$

$W > 0$ establishes that the frozen writer produced different persistent states. It does *not* establish that every textual atom is a pure causal effect of the reward bit, because decoding randomness is not separately noise-calibrated at the atom level. We therefore use the same-mode wording control in Section 4 to test the stronger alternative that ordinary prompt wording alone explains the observed write divergence.

2.2 NATIVE TRANSPORT IS A SEQUENCE OF DISTINCT OBSERVABLES

For a held-out future query q , define native retrieval exposure at the source-item level,

$$E_r(q) = \mathbf{1}[m^{(r)} \in \text{Retrieve}(q; \mathcal{M}^{(r)})].$$

Exposure is an availability variable: it says that the branch-conditioned source-memory item reached the policy context. It does not identify which treatment-specific content inside that item was read or used by the policy. In particular, source-item exposure is weaker than *treatment-residual exposure*: the latter would require showing that the branch-differentiating content, rather than merely the shared source item, is represented in the policy’s effective readout. The frozen logs support the former but do not identify the latter.

We therefore measure the first downstream decision boundary separately. Let $p_r^{(1)}(\cdot \mid q)$ be the empirical distribution over the normalized first structured action emitted under branch r for the same matched future-task state. The support is the set of first structured actions observed in either branch for that state; actions absent from one branch receive zero mass there. The action-level contrast is

$$U(q) = \text{TV}(p_S^{(1)}(\cdot \mid q), p_F^{(1)}(\cdot \mid q)).$$

Finally, for terminal outcome $y_r(q)$,

$$O(q) = |y_S(q) - y_F(q)|.$$

A *branch-comparison state* throughout the paper means one matched future-task state evaluated under the paired success- versus failure-conditioned memory branches while the non-memory task state is held fixed. The reported 36 Shopping states are 36 such matched comparison units; they are not 36 independent source trajectories.

The native scientific object is therefore the ordered observation tuple

$$\mathcal{T}_{\text{native}}(q) = (W, E, U, O),$$

not any one component in isolation.

Crucially, these components are **not commensurate**. W is a memory-state distance, E is an exposure indicator/rate, U is a distributional action distance, and O is an endpoint contrast. We therefore do not divide them into a scalar “transport efficiency” or treat their numerical differences as a mediation decomposition. The order is mechanistic; the scales are not.

A further identifiability issue is *evaluation coarsening*. Let $\Pi_A(\mathcal{T}_{\text{native}})$ project the chain onto observed stages $A \subseteq \{W, E, U, O\}$. The projection is many-to-one: for example, a small endpoint under $\Pi_{\{O\}}$ is compatible with writer failure, retrieval absence, post-retrieval filtering, or sparse heterogeneous transport, while $\Pi_{\{E\}}$ cannot distinguish availability from policy uptake. We therefore retain the categorical stage signature $\Sigma = (z_W, z_E, z_U, z_O; z_F)$, with z_F denoting the forced-capacity side control. It is not a scalar metric; it preserves which explanations remain distinguishable after each stage is observed.

2.3 FORCED LEVERAGE IS A CAPACITY CONTROL, NOT A NATIVE STAGE

The forced fixed-evidence experiment mechanically supplies paired memory content to the downstream task packet. Let

$$L_{\text{forced}}(q) = |y^{\text{forced},S}(q) - y^{\text{forced},F}(q)|.$$

Because native retrieval is bypassed, L_{forced} is not inserted between W and E . It answers a different question: *can* branch-specific memory content affect the downstream system when supplied? Native (E, U, O) asks whether the deployed memory path *does* transport that contrast. This capacity-versus-transport separation is the key control that prevents a weak native endpoint from being summarized as “memory can never matter.”

2.4 AN ORDINAL STAGE-EVIDENCE LADDER

The central measurement problem is that W , E , U , and O cannot be compared on one numerical scale, and their frozen tests do not have matched sample units or statistical power. We therefore localize an *evidence boundary* with an ordinal ladder rather than a post-hoc transport coefficient or a latent bottleneck estimator. Let the ordered native stages be

$$s_1 = \text{write}, \quad s_2 = \text{retrieval exposure}, \quad s_3 = \text{first-action uptake}, \quad s_4 = \text{terminal outcome}.$$

For each stage we preserve its native evidence semantics. A stage may be SUPPORTED by its frozen inferential test, OBSERVED directly when the claim is descriptive rather than thresholded, or NOT-SUPPORTED when its frozen primary test does not establish a stable branch contrast. We then define

$$b^* = \min\{j : s_j \text{ is not supported/observed while } s_1, \dots, s_{j-1} \text{ are supported/observed}\}.$$

This is an ordering rule, not a ratio: no cross-stage units are divided, normalized, or interpreted as mediation mass. Because the stage tests are not power-matched, b^* means “first unsupported measured stage under the frozen tests,” not “the stage where the true causal effect first becomes small.”

The forced fixed-evidence arm is evaluated separately as a *side capacity condition*. It can weaken the global-insensitivity explanation, but because it bypasses native retrieval it cannot make an otherwise unsupported native stage pass. This prevents the positive forced endpoint from being inserted into the native chain merely because it is statistically strong.

Multiplicity and interpretation. The frozen tests answer different stage-specific questions and are not searched as a family for the smallest p -value. We therefore do not make an omnibus “at least one stage is significant” discovery claim and do not use a post-hoc multiplicity correction to choose b^* . The boundary operator depends on the preregistered ordering and each stage’s own frozen evidence rule. Reported p -values should consequently be read as stage-local evidence summaries, not as interchangeable entries in a single multiple-testing family.

The same ladder can be read as an alternative-explanation audit:

A1: no state intervention $\Rightarrow W \approx 0$,

A2: global downstream insensitivity $\Rightarrow L_{\text{forced}} \approx 0$,

A3: retrieval absence only $\Rightarrow E \approx 0$,

A4: exposure is uptake $\Rightarrow$ observed E should be accompanied by supported U ,

A5: universal directional transport $\Rightarrow O$ should be broadly nonzero with stable sign.

In our frozen evidence, write is SUPPORTED, native source-item exposure is OBSERVED, first-action uptake is NOT-SUPPORTED, and terminal transport is sparse and sign-heterogeneous. Hence $b^* = s_3$: **first-action uptake is the first unsupported measured native stage after write and source-item exposure have been established.** The positive forced side control further weakens a global-insensitivity account without changing b^* ; it does not identify a native latent attenuation onset.

This is still not an identified mediator coefficient. A causal mediation statement would additionally require compatible intervention surfaces and a noise/seed design that isolates each branch-specific atom. We do not have that evidence and do not make that claim.

2.5 INTERPRETATION DISCIPLINE

For clarity, we use four terms narrowly: **write** means that the persistent state changes; **exposure** means that the branch-conditioned source item appears in native retrieved context; **uptake** means that a measured policy decision distribution changes across branches; and **capacity** means that supplied branch-specific memory can affect downstream behavior when native retrieval is bypassed. These are measurements on different surfaces, not synonyms for “memory use.”

Three non-implications are load-bearing throughout the paper:

persistent-state divergence $\nRightarrow$ behavioral divergence,

retrieval exposure $\nRightarrow$ policy uptake,

forced leverage $\nRightarrow$ native end-to-end transport.

These distinctions also explain why our stopped evidence-gated residual extension is not a negative method result: that extension failed qualification before a behavioral method-effect experiment was authorized. The present paper therefore remains an identification/measurement study of the transport chain.

3 STAGE 1: REWARD-CONDITIONED WRITING ROBUSTLY CHANGES PERSISTENT STATE

We first test the write stage with paired interventions over released WebArena Shopping trajectories. Within each pair, the task text and full action history are fixed before the writer is called. One branch uses the released success-reflection objective and the other uses the released failure-reflection objective; the resolved writer model is matched within pair and temperature is zero. Raw responses and the pre-writer trajectory projection are content-addressed before analysis.

The initial F0 tranche requested six source trajectories. Four returned complete memory outputs in both conditions; the two incomplete requests terminated in the failure-label writer before assistant text existed and were excluded from the paired content estimand. Among the four complete pairs, every pair changed normalized memory content and the extracted memory-item title set, with mean token-set Jaccard distance 0.735. Receipt-level forensics preserve the two incomplete arms as provider-response failures rather than treating them as negative memory examples.

The later frozen breadth tranche expands the write-stage evidence to 20 complete Shopping success/failure pairs. All 20/20 pairs produce different branch memories, with pooled token-set Jaccard distance 0.673. When combined with the four Reddit pairs used in the cross-domain replication, the D0 lineage contains 24/24 paired source experiences with the same pre-writer trajectory projection, the same resolved writer model within pair, temperature zero, and different branch-memory content. This establishes a robust operational branch contrast at the persistent-state level.

The content changes are not only token-level. A deterministic post-ready strategy-slot audit maps frozen memory titles into navigation, pagination/coverage, query expansion, semantic matching, evidence capture, extraction, and verification slots. In the original four complete F0 pairs, the strategy-slot set changes in all four. We treat this as a structural diagnostic rather than a semantic-validity claim: the current data do not certify each textual delta atom as a pure reward-caused proposition because explicit decoding seeds or same-condition noise-floor replications are unavailable.

The write-stage result therefore supports a narrow statement: under the released writer policy, flipping the reward-conditioned reflection branch on the same source experience reliably produces a different durable state. It does not by itself establish that the difference will be retrieved, followed, or reflected in terminal behavior. Section 4 first tests the strongest simple prompt-sensitivity reduction; Section 5 then measures downstream transport stage by stage.

4 RULING OUT GENERIC PROMPT SENSITIVITY

The success/failure writer-branch intervention changes the reflection instruction together with the outcome/reward semantics. It is therefore a bundled protocol contrast, not a pure reward-bit intervention. A generic prompt-sensitivity account predicts that a sufficiently strong wording perturbation inside one reflection mode should produce comparable memory divergence; the following control tests that reduction, but cannot decompose instruction semantics from reward semantics within the released branch. We evaluate it on eight released trajectories disjoint from F0 and balanced across original outcomes.

Before provider calls, we freeze semantic-preserving rewordings for both reflection modes. Under token-set Jaccard distance, the released success-versus-failure prompts differ by 0.114. The within-success rewording differs by 0.171 and the within-failure rewording by 0.152. The control thus perturbs surface wording more strongly while preserving the reflection objective.

Each fresh trajectory is written with the released success mode and the released failure mode. The same writer also processes one frozen semantic-preserving rewording of each mode at temperature zero. Let B_i be the memory distance between the two released reward modes and let W_i be the mean within-mode rewording distance. The frozen paired statistic is

$$\Delta_{\text{prompt}} = \frac{1}{8} \sum_{i=1}^8 (B_i - W_i). \quad (1)$$

All 32 control outputs complete. Mean reward-mode distance is 0.700, while the stronger same-mode wording control has mean 0.595. The paired excess is 0.105, and an exact one-sided sign-flip test gives $p = 0.0078$. This passes the frozen 0.10 effect margin together with the $p < 0.05$ gate. The reward-conditioned write effect therefore survives a same-information prompt-wording reduction that is deliberately stronger in lexical perturbation.

5 STAGE-RESOLVED TRANSPORT EVIDENCE

The write intervention becomes scientifically useful only once alternative explanations for weak downstream transport evidence are separated. We therefore use the frozen evidence in three roles: forced fixed-evidence exposure as a *capacity control*, native retrieval as an *availability measurement*, and first-action/terminal contrasts as *native transport measurements*. This section shows that their joint pattern is more informative than any endpoint alone.




Figure 1: Write-channel evidence and its strongest simple control. (a) Among the four F0 pairs complete in both writer conditions, flipping only the reward-conditioned writer mode changes every memory. (b) On eight fresh control trajectories, reward-mode memory divergence exceeds stronger semantic-preserving same-mode prompt rewording. The control mean is 0.595 versus 0.700 between reward modes; paired excess is 0.105 with exact $p = 0.0078$.

5.1 FORCED INJECTION ESTABLISHES DOWNSTREAM CAPACITY, NOT NATIVE TRANSPORT

We first bypass native retrieval and directly inject each paired memory into a fixed future-task packet. This asks whether the downstream system can respond to the paired persistent states when the experimenter supplies them.

The original fixed-state probe contains 12 aligned success-memory versus failure-memory rollouts; 7/12 choose different structured next actions. We then freeze all four complete initial F0 source-memory pairs and four future tasks before terminal-policy execution. Every source pair is crossed with every future task, giving 16 cells. The confirmatory run uses eight fresh rollouts per memory condition in every cell, for 256 completed calls. Mean absolute terminal success-rate difference is 0.15625, above the frozen 0.15 practical floor, and the 100,000-draw within-cell permutation test gives $p = 0.00074$.

This result weakens the explanation that the downstream system is simply incapable of responding to memory content. It does *not* estimate native transport because exposure is mechanically supplied. The cell effects are also heterogeneous: eight cells are zero, six have positive failure-minus-success effects, and two have negative effects; the two largest cells account for 78.2% of squared-effect mass. Even under forced exposure, a universal harmful direction is unsupported.

5.2 NATIVE SOURCE-ITEM EXPOSURE REMAINS COMMON WHILE ACTION UPTAKE IS UNSUPPORTED

We next restore the native memory-bank pathway. After expanding the frozen Shopping memory bank, the branch-conditioned source-memory item is retrieved on 125 of 172 held-out retrieval opportunities, a rate of 0.727. We intentionally call this *source-item exposure*: the retrieval log proves that the paired source item entered native context, but it does not prove that the policy attended to or represented the branch-differentiating residual within that item. Establishing such *treatment-residual exposure* would require a finer intervention, such as residual masking, diff-aware attribution, or a matched readout probe, none of which is present in the frozen evidence. Thus the 125/172 result rules out “the relevant source item was never retrieved” as the sole account, but it does not certify causal consumption of the treatment-specific content. In particular, the shortcut “Retrieval hit is a surrogate for policy use” is not valid for this evaluation.

Despite that exposure, the first measured decision boundary is weak. Each of the frozen 36 Shopping branch-comparison states is one matched future-task state evaluated under the paired success- and failure-conditioned memory branches; the first-action distribution is formed over normalized first structured actions observed across the two branches for that same state. Mean total-variation distance

Table 1: Stage-evidence ladder. Evidence states preserve the semantics of each measurement instead of forcing heterogeneous quantities onto one scale. The forced arm is a side capacity control and therefore does not enter the native-stage ordering.

Stage	Evidence state	Frozen observation / interpretation
Persistent write	Supported	20/20 Shopping pairs diverge; stronger same-mode control leaves excess 0.105, exact $p = 0.0078$.
Forced capacity (side control)	Supported	$ \Delta = 0.15625$, $p = 0.00074$; supplied memory has downstream leverage but native retrieval is bypassed.
Native source-item exposure	Directly observed	125/172 held-out opportunities retrieve the branch-conditioned source-memory item; this is item availability, not treatment-residual exposure or policy use.
First-action uptake	Not supported	$TV = 0.06944$, $p = 0.5801$, 0/36 modal changes over matched first structured-action distributions; this is the first unsupported <i>measured</i> native stage after observed prerequisites.
Native outcome	Sparse / heterogeneous	Shopping 34/36 zero; Reddit 6/8 zero and the two nonzero signs are opposite.

is 0.06944 with $p = 0.5801$, and 0/36 states change their modal first action. This probe is deliberately narrow: it tests the earliest structured decision that is consistently available in the frozen logs and does not rule out a later second-step, tool-sequence, or plan-level difference. A post-hoc working-memory diagnostic yields branch shift 0.00335 with $p = 0.2052$. Native terminal transport is also small: mean absolute success difference is 0.02083 with permutation $p = 0.4289$, and 34/36 cells are exactly zero.

The conjunction is the stage-resolved evidence result: the strongest supported operational localization lies after exposure and before stable action uptake. Applying the ordinal rule from Section 2.4, write is SUPPORTED, native source-item exposure is OBSERVED, and first-action uptake is the first *measured* native stage whose frozen primary test does not support a stable branch contrast. Thus $b^* = \text{first-action uptake}$ is an *evidence-localization boundary*. It is not a claim that first action is the causal onset of attenuation, nor that no later policy surface can differ. Terminal sparsity is downstream corroboration rather than the quantity used to choose the boundary. Because the stage tests differ in units and statistical sensitivity, this does not identify where a latent causal effect physically attenuates; it is not a causal mediation coefficient.

A structural sensitivity audit clarifies what is load-bearing in this localization. Removing terminal outcome evidence, removing the Reddit terminal tranche, or removing the forced-capacity side control leaves the first unsupported measured native stage at first-action uptake. In contrast, merging native exposure with first-action uptake removes the diagnostic resolution needed to distinguish availability from behavioral use. The contribution is therefore not simply to report more stage metrics: for this frozen system, the *separation of exposure from uptake* is the load-bearing measurement that prevents these two explanations from becoming observationally aliased.

The full six-hypothesis alternative-explanation audit is reported in Appendix Table 5; the main-text stage ladder retains the decision-critical evidence states without spending a second full table on the same inference chain.

5.3 CROSS-DOMAIN BOUNDARY: WRITE REPLICATION WITHOUT MONOTONE TRANSPORT

We repeat the write-versus-native-transport measurement on Reddit. Reward-conditioned writing again diverges in 4/4 paired source memories. Native terminal transport is larger than Shopping in magnitude but remains unresolved and highly sparse: mean absolute success difference is 0.125 with $p = 0.2253$, 6/8 cells are zero, and the two nonzero cells have opposite signs. The second domain therefore supports the same coarse structural separation—robust write divergence need not become stable native terminal transport—while rejecting a simple universal attenuation constant or signed harm coefficient. Reddit does not contain the matched native exposure and first-action measurements needed to replicate the Shopping evidence boundary.




Figure 2: **Native transport and forced capacity are different experimental surfaces.** The horizontal path is the deployed reuse chain: persistent write $\rightarrow$ native source-item exposure $\rightarrow$ first-action uptake $\rightarrow$ native outcome. The forced fixed-evidence arm bypasses retrieval and is shown as a side capacity control. On the frozen Shopping tests, first-action uptake is the first unsupported measured native stage after write and item exposure are established; this is an evidence boundary, not a latent causal bottleneck or mediation coefficient.

An outcome-blind structured-memory control gives a complementary warning about endpoint interpretation. It produces a small terminal contrast of 0.04514 ($p = 0.0048$), below the frozen 0.15 practical floor, with 32/36 cells zero. Statistical detectability alone is therefore not the scientific target; a sparse small effect can pass a null test without matching the magnitude or mechanism of forced branch leverage.

5.4 WHY THE ONE-STEP REWARD-TO-BEHAVIOR STORY FAILS

A one-step interpretation would take the 256 forced rollouts as evidence that “reward error propagates to future behavior.” The native measurements show why this is too coarse. Forced injection estimates leverage conditional on making a particular memory state available; native transport additionally depends on retrieval competition, context composition, and policy uptake. Indeed, coarsened views of the same frozen system disagree: write-only looks strongly affected, retrieval-only looks active, native endpoint-only looks nearly null, and forced-only shows leverage. The full stage signature reconciles these summaries rather than averaging them. The difference between 0.15625 forced leverage and 0.02083 native Shopping terminal transport is therefore not a scalar attenuation factor to divide or average; the two surfaces answer different questions.

Conversely, looking only at the native endpoint would also lose the mechanism: a 0.02083 Shopping terminal difference does not tell us that the writer failed, that memory was unavailable, or that the downstream system was globally insensitive. The stage-resolved controls exclude or weaken each of those simpler accounts before the first-action boundary is reached.

The evaluation consequence is concrete. For a persistent update, report at least one state-change metric, one exposure measurement, and one downstream uptake/outcome metric; use a capacity control only if its bypassed stages are explicit. Where possible, distinguish source-item exposure from treatment-residual exposure rather than treating retrieval of a shared carrier item as proof that the causal residual was consumed. Then report the first measured stage where evidence for transport is no longer supported, while keeping stage-specific power and measurement semantics explicit, rather than crediting the update with behavioral authority because its representation changed or because a retrieved item appeared in context.

6 RELATED WORK

Feedback-conditioned memory. Prior work covers reward reliability (Wang et al., 2026), success/failure-derived memory (Ouyang et al., 2026), reward-conditioned construction (Wang

et al., 2025), and scoring/reuse distortion (Asadolahi et al., 2026; Yang et al., 2026). These are not our novelty; the frozen writer branch is only our controlled intervention.

Retrieval, utilization, and our residual. Mem2ActBench studies active use (Shen et al., 2026); QCR separates retrieval from reuse (Li et al., 2026); Yuan et al. diagnose write/retrieval/utilization (Yuan et al., 2026); and lifecycle work names the same broad stages (Lin et al., 2026). These concepts are not our novelty. Our residual question is intervention-specific: after the same source trajectory is converted into two branch-conditioned persistent states, which measured boundary still carries that branch contrast? The load-bearing distinction is separation of exposure from uptake while preserving the identity of the controlled branch intervention, not lifecycle naming or a power-matched causal bottleneck. A retrieval hit can establish source-item availability without establishing transport of the branch-differentiating residual into the measured decision. Known self-improvement variance provides complementary motivation (Ye et al., 2026).

7 LIMITATIONS

The released success/failure branch jointly changes reflection instruction and outcome/reward semantics; decoding-seed binding and same-condition noise-floor replication are unavailable, two failure-label calls ended before assistant text, and source-faithful live WebArena navigation is unavailable. We therefore estimate a bundled writer-protocol contrast, not a pure reward-bit effect.

Stage tests differ in units, samples, power, and intervention surfaces, so first-action uptake is only an *evidence-localization boundary*, not an identified latent attenuation onset or mediator. The 125/172 statistic is source-item retrieval, not treatment-residual exposure or policy use. The frozen logs do not identify which branch-differentiating content is actually consumed downstream. Forced injection bypasses native retrieval/context composition and measures supplied-evidence capacity.

Shopping is sparse (34/36 zero); Reddit is sparse, sign-heterogeneous (6/8 zero), and lacks matched exposure/action measurements, so neither a universal effect nor cross-domain action boundary is claimed. External validity is restricted to the frozen released writer/policy configuration and the Shopping+Reddit substrate evaluated here; other writers, backbones, domains, and memory architectures are untested. The CBRG extension stopped before behavioral evaluation; no behavioral repair efficacy is claimed.

8 CONCLUSION

The same feedback-conditioned writer intervention reveals two distinct facts about persistent self-improvement. First, changing the released success/failure reflection branch on the same source experience reliably writes different durable memory. Second, that written difference does not automatically become a native behavioral difference: on Shopping, source-item exposure remains common while first-action uptake is the first unsupported measured native stage, and terminal transport is sparse. Forced exposure shows that the branch-specific memories can influence behavior when supplied, so the native attenuation cannot be reduced to a claim that the downstream policy is globally insensitive to memory. The resulting lesson is two-sided: evaluative feedback can become persistent agent state, but **state divergence is not behavioral authority**. A sound evaluation must therefore ask both how differences are written and whether the deployed reuse chain actually transports them into decisions.

REFERENCES

- Mohammad Asadolahi, Amir Amini, Samira Talebi, Amirfarhad Farhadi, and Azadeh Zamanifar. Memory reward inflation in self-improving llm agents, 2026.
- Yifei Li, Heng Wang, Lingling Zhang, Muye Huang, Xinyu Zhang, Jiashuai Liu, Hang Yan, and Rongman Xu. Beyond retrieval: Query-conditioned reuse of long-horizon agent trajectories, 2026.
- Zehao Lin, Xixuan Hao, Renyu Fu, Shaobo Cui, Kai Chen, Chunyu Li, Zhiyu Li, and Feiyu Xiong. A survey on long-term memory security in llm agents: Attacks, defenses, and governance across the memory lifecycle, 2026.

Siru Ouyang, Jun Yan, I-Hung Hsu, Yanfei Chen, Ke Jiang, Zifeng Wang, Rujun Han, Long T. Le, Samira Daruki, Xiangru Tang, Vishy Tirumalashetty, George Lee, Mahsan Rofouei, Hangfei Lin, Jiawei Han, Chen-Yu Lee, and Tomas Pfister. Reasoningbank: Scaling agent self-evolving with reasoning memory. In *International Conference on Learning Representations*, 2026.

Yiting Shen, Kun Li, Wei Zhou, and Songlin Hu. Mem2actbench: A benchmark for evaluating long-term memory utilization in task-oriented autonomous agents, 2026.

Jiaxuan Wang, Yulan Hu, Wenjin Yang, Zheng Pan, Xin Li, and Lan-Zhe Guo. Aligning agents via planning: A benchmark for trajectory-level reward modeling, 2026.

Yu Wang, Ryuichi Takanobu, Zhiqi Liang, Yuzhen Mao, Yuanzhe Hu, Julian McAuley, and Xiaojian Wu. Mem- α : Learning memory construction via reinforcement learning, 2025.

Yi Yang, Zhennan Chen, Yihong Zhuang, et al. Romerl: Balancing feedback coverage and the memory-reward trap in self-evolving agent memory via reduced-order utility states, 2026.

Qinyuan Ye, Yu Li, Yada Pruksachatkun, Jiaxin Zhang, and Chien-Sheng Wu. On the fragility of self-improving agents: Variance, task order, and underspecification, 2026.

Boqin Yuan, Yue Su, and Kun Yao. Diagnosing retrieval vs. utilization bottlenecks in llm agent memory, 2026.

A F0 REPRODUCTION DETAILS

A.1 RELEASED TRAJECTORY SAMPLE

The F0 sample uses six task IDs from the released AWM shuffle-seed-42 Shopping parquet: 21 through 25, plus 47. The deterministic sampler selects the lowest task IDs with parseable trajectory records within each original-outcome stratum. Original failures are 21 through 22 plus 24. Original successes are 23 plus 25 and 47.

A.2 MEMORY CONSTRUCTION INTERVENTION

Each released trajectory is summarized into an action trace. The same trace is inserted into two prompts. One prompt contains the released ReasoningBank success-reflection system instruction. The paired prompt contains the released failure-reflection system instruction. The task prompt stays fixed. The writer model stays fixed. Temperature stays at zero. The raw provider output is archived before metric computation.

A.3 PAIR-LEVEL F0 RESULTS

Table 2: Complete paired interventions.

Task	Original outcome	Content changed	Token distance
21	failure	yes	0.691489
22	failure	yes	0.708108
23	success	yes	0.769953
25	success	yes	0.769608

The extracted memory-item title set changes for every complete pair. The mean token-set Jaccard distance is 0.734789. A deterministic post-ready strategy-slot audit maps the frozen titles into navigation, pagination/coverage, query expansion, semantic matching, evidence capture, extraction, and verification/completeness. All four complete pairs change their slot set. The slot-set Jaccard distances are 0.571, 0.400, 0.500, and 0.500 for tasks 21, 22, 23, and 25, respectively, with mean 0.493. This is a structural diagnostic over frozen titles, not an embedding-based semantic claim.

Provider failures for task 24 and task 47 are excluded from the paired analysis. Both are failure-label calls with `ArkResponseStateError`; neither produces assistant text or a raw memory SHA.

Private response identifiers are retained and GET-only recovery finds no text. The four completed failure-label calls produce 777, 852, 1,113, and 1,296 output tokens. These receipts distinguish provider response-state failures from a later memory-parser rejection, while leaving arm-dependent completion bias unresolved.

B PROMPT-PERTURBATION CONTROL DETAILS

The control set is disjoint from the F0 requests. Each fresh trajectory is written under both released reward modes. One semantic-preserving rewording is also evaluated within each mode. Table 3 reports the paired distances used in the frozen sign-flip test.

Table 3: Full fresh prompt-perturbation control. Between is the released success-versus-failure memory distance; Within averages the two same-mode rewording distances.

Task	Original outcome	Between	Within	$B_i - W_i$
26	failure	0.726	0.662	0.065
50	failure	0.668	0.599	0.069
51	failure	0.776	0.534	0.242
96	failure	0.682	0.549	0.134
48	success	0.716	0.656	0.059
49	success	0.573	0.594	-0.022
117	success	0.703	0.593	0.110
126	success	0.755	0.572	0.182
Mean		0.700	0.595	0.105

C DOWNSTREAM CONFIRMATORY DETAILS

C.1 FIXED-STATE ACTION DISTRIBUTIONS

The distributional action stress test retains all four intervention states from the action-reach experiment and uses eight rollouts per memory condition. The per-state success-memory versus failure-memory TV distances are 0.625 for task 22 and 0.375 for task 23. They are 0.125 for task 24 and 0.250 for task 26. Their mean is 0.34375. The frozen 100,000-draw permutation audit gives $p = 0.310527$.

C.2 FROZEN TERMINAL MATRIX

The terminal confirmatory experiment reuses every complete F0 pair and every future task frozen for the earlier fixed-evidence run. Source-memory tasks form $\{21, 22, 23, 25\}$. Future tasks form $\{164, 385, 387, 388\}$. The design therefore contains 16 source-future cells. Each cell receives eight fresh success-memory rollouts and eight fresh failure-memory rollouts. The policy model is Doubao Seed 2.0 Mini at temperature 0.2. All 256 requested calls complete with zero provider failures.

The primary statistic is the mean across cells of the absolute difference between success-memory and failure-memory terminal benchmark success rates. The support gate is frozen before calls at an effect floor of 0.15 together with a within-cell permutation $p < 0.05$. The permutation audit uses 100,000 draws with seed 20260822. The observed statistic is 0.15625 and the audit gives $p = 0.00074$.

C.3 HETEROGENEITY AND CORRUPTION-RATE DECOMPOSITION

For each frozen cell, the confirmatory run estimates paired conditional success probabilities. Eight of the 16 cells have zero observed success-rate contrast. Six have positive failure-memory minus success-memory effects and two have negative effects. The largest squared effects occur at source/future pairs 22/388 and 25/387; together they contribute 78.2051% of the total squared-effect mass. This sign and magnitude heterogeneity is why the revised paper treats the experiment as a forced-leverage measurement rather than a uniform directional effect or a native end-to-end transport estimate.

If reward-label corruption selects the failure-conditioned memory with probability q , the induced between-memory variance equals $q(1-q)\Delta_c^2$ in cell c . The exact mean squared conditional effect across all 16 cells is 0.076172. Averaging over cells gives the curve $0.076172q(1-q)$, with value

Table 4: F2R1 confirmatory terminal success rates. S denotes success-label memory and F denotes failure-label memory.

Source	Future	S rate	F rate	$ \Delta $
21	164	1.000	1.000	0.000
21	385	0.000	0.250	0.250
21	387	0.125	0.250	0.125
21	388	1.000	0.625	0.375
22	164	1.000	1.000	0.000
22	385	0.000	0.125	0.125
22	387	0.125	0.000	0.125
22	388	0.250	1.000	0.750
23	164	1.000	1.000	0.000
23	385	0.000	0.000	0.000
23	387	0.125	0.250	0.125
23	388	1.000	1.000	0.000
25	164	1.000	1.000	0.000
25	385	0.000	0.000	0.000
25	387	0.125	0.750	0.625
25	388	1.000	1.000	0.000
Mean absolute difference				0.15625

0.006855 at $q = 0.1$, value 0.014282 at $q = 0.25$, and value 0.019043 at $q = 0.5$. This remains a descriptive plug-in calculation conditional on forced memory exposure; it is not used as the main claim in the revised stage-resolved paper.

C.4 NATIVE TRANSPORT AUDIT ADDED AFTER THE FORCED-LEVERAGE STUDY

The later frozen native-pipeline audit changes the interpretation of the forced matrix. In Shopping, the branch-conditioned source-memory item is retrieved in 125/172 held-out opportunities after bank expansion, but the mean absolute native terminal success difference is only 0.02083 with permutation $p = 0.4289$; 34/36 terminal cells are zero. The corresponding first-action total-variation statistic is 0.06944 with $p = 0.5801$, and no frozen state changes its modal first action between reward branches. A post-hoc working-memory branch-shift diagnostic is 0.00335 with $p = 0.2052$.

The cross-domain Reddit audit contains four complete success/failure writes, all of which diverge. Its native terminal mean absolute difference is 0.125 with $p = 0.2253$; 6/8 cells are zero and the two nonzero cells have opposite signs. These later measurements are why the main text separates forced capacity from native transport and treats first-action uptake as the first unsupported measured Shopping stage after write and source-item exposure, rather than as an identified latent attenuation onset or causal mediation result.

For finite-cell sensitivity, we run a fixed-seed nonparametric bootstrap with 100,000 repetitions, resampling 16 source-future cells with replacement. The percentile 95% interval for mean absolute effect is [0.054688, 0.273438]. The corresponding interval for mean squared effect is [0.010742, 0.164062], which maps to [0.002686, 0.041016] for $V(0.5)$. This bootstrap is descriptive sensitivity to the observed cell composition. The preregistered within-cell permutation test on the 256 confirmatory rollouts supplies the primary significance result.

C.5 FULL ALTERNATIVE-EXPLANATION AUDIT

Table 5: Alternative-explanation audit. “Weakened” means the frozen evidence makes the explanation incomplete but does not mathematically eliminate every variant of it.

Explanation	Frozen evidence	Disposition
No persistent state intervention	20/20 Shopping and 4/4 Reddit writes diverge; same-mode excess = 0.105, $p = 0.0078$	Inconsistent with the observed write-stage intervention.
Downstream policy is globally insensitive to memory	Forced terminal $ \Delta = 0.15625$, $p = 0.00074$	Weakened: supplied branch-specific memory has measurable leverage.
Native null is only retrieval absence	125/172 native Shopping source-item retrieval hits	Weakened: substantial item exposure coexists with weak first-action contrast.
Retrieval hit is a surrogate for policy use	first-action TV = 0.06944, $p = 0.5801$; 0/36 modal changes	Rejected as an evaluation equivalence: item availability is not policy uptake.
Universal directional terminal transport	Shopping 34/36 zero; Reddit 6/8 zero, two nonzero signs opposite	Not supported by the frozen native cells.
Causal mediation has been identified	no per-atom seed binding/noise-floor replication; forced/native surfaces differ	Unresolved and explicitly outside the claim set.