

DO TEMPORAL SKILLS REALLY REPAIR AGENTS? AN INTERVENTION AUDIT OF REPAIR AND ATTRIBUTION

Anonymous authors

Paper under double-blind review

ABSTRACT

When an agent repeatedly fails on evolving data, a reusable “skill” can appear to fix the problem even when the comparison itself is misleading. We study how to audit that claim rather than how to learn another skill. Our intervention ladder separates net repair relative to the original agent ($T-N$), a same-surface operation-output contrast ($T-G_0$), the placebo/surface perturbation itself (G_0-N), and the effect of exposing the exact same operation output through callable state rather than ordinary context ($T-R_{\text{surf}}$). We do not collapse $T-G_0$ into a pure operation effect when G_0 differs from N . This distinction changes scientific verdicts. On one DeepSeek cutoff cell, $T=100\%$ and $G=72\%$ but $N=100\%$, so the apparent +28-point two-arm gain is no repair. Fresh DeepSeek grounding remains directionally positive under the stronger controls (+30 and +20 points for $T-N$ and $T-G_0$ in two five-endpoint cells with zero observed G_0-N shift), but those cells are low-resolution. A post-hoc EIA compatibility probe illustrates why the decomposition matters: $T-N$ is +37.5 points while G_0-N is -25 points, so we classify it as net repair with surface interaction rather than confirmation. A Kimi grounding support cell is also downgraded. Finally, R_{surf} materializes the byte-identical targeted-operation output in context. Across 18 pre-frozen DeepSeek endpoints $T-R_{\text{surf}}$ is 0.0 points; the paired TOST passes at the prospectively frozen ± 10 -point portfolio margin, with paired-t 95% CI $[-8.5, 8.5]$ and bootstrap 95% sensitivity CI $[-8.3, 8.3]$. Because an observed-variance sensitivity calculation is only about 53%, the portfolio is ceiling-heavy, and the strictly non-ceiling subset is unresolved, we treat this only as a low-sensitivity frozen-margin surface check. Thus stronger controls can preserve bounded net repair, expose surface interaction, reverse support cells, and remove skill-container credit within the tested harness without implying broad equivalence to temporal-retrieval algorithms. The contribution is the audit and the verdict changes it forces.

1 INTRODUCTION

Suppose an agent is asked: *What was the latest value known by the end of April?* A May release is topically relevant, easy to retrieve, and wrong for the question. On the next period the numbers change, but the required operation does not: filter evidence by the information cutoff. Similar recurring operations arise when an agent must distinguish publication time from reporting time or ground a numerical change in documentary context available at the time. Reusable skills are an appealing way to persist such procedures, but a successful skill run does not by itself identify what was repaired or what part of the implementation deserves credit.

The first problem is the comparator. A targeted helper can beat a generic helper simply because the generic helper degrades the original agent. We observe exactly this pattern on a frozen DeepSeek cutoff cell: targeted T succeeds on 100%, target-blind G on 72%, and the original agent N on 100% of endpoint-repeat observations. A conventional T -versus- G evaluation reports +28 points; adding N changes the verdict to *no repair*. This is not a variance correction. It changes which scientific claim is admissible.

Figure 1 turns this identification problem into the running example for the paper. The same cutoff question can make T look beneficial against a mechanism-misaligned stress helper while showing no improvement over the original agent. The lower row previews the next attribution problem: even




Figure 1: **Scene: a targeted skill can look helpful because its comparator is harmful.** On one frozen DeepSeek cutoff cell, T and the original agent N are both 100%, whereas the target-blind stress helper G is 72%. A T-versus-G comparison therefore shows +28 points even though T–N is exactly zero and no repair occurred. The bottom row previews the attribution controls used later: G_0 isolates same-surface no-operation exposure, and R_{surf} holds the exact T output fixed while changing only its integration surface.

when net repair exists, helper exposure, targeted operation output, and where that output enters the harness are distinct interventions and need distinct controls.

The second problem is attribution. We therefore keep four observable contrasts separate:

$$\underbrace{\Delta_{T-N}}_{\text{net repair}}, \quad \underbrace{\Delta_{T-G_0}}_{\text{same-surface output contrast}}, \quad \underbrace{\Delta_{G_0-N}}_{\text{surface/placebo perturbation}}, \quad \underbrace{\Delta_{T-R_{\text{surf}}}}_{\text{integration-surface effect}}. \quad (1)$$

G_0 is an information-empty no-operation placebo exposed through the same helper surface as T. $T-G_0$ therefore asks whether the targeted operation output changes behavior *conditional on that surface*; G_0-N separately reveals whether the surface/placebo itself perturbs the original agent. We do not call $T-G_0$ a pure operation effect when G_0-N is nonzero. Algebraically, $T-N=(T-G_0)+(G_0-N)$, so the three contrasts expose rather than hide a surface interaction. R_{surf} asks a different question: it runs the identical frozen operation on identical candidate evidence but materializes the exact output in ordinary context instead of callable state. It is deliberately a controlled surface intervention, not an independently learned temporal retriever.

This audit changes several interpretations. Fresh DeepSeek grounding runs have +30 and +20 point $T-N$ and $T-G_0$ directions in C3/C4 with zero observed G_0-N shift on those runs, although each cell has only five endpoints. In a post-hoc EIA compatibility probe, $T-N$ is +37.5 points but G_0-N is –25 points; we use this cell primarily to illustrate net repair with surface interaction, not as confirmation. A Kimi same-domain grounding support cell is downgraded under the fresh controls. The surface audit narrows credit again: across 18 pre-frozen DeepSeek load-bearing endpoints $T-R_{\text{surf}}$ averages 0.0 points. A formal two-one-sided equivalence test at the prospectively frozen ± 10 -point portfolio margin passes; paired-t and bootstrap 95% intervals are $[-8.5, 8.5]$ and $[-8.3, 8.3]$ points, respectively. Because a post-hoc observed-variance sensitivity calculation is only about 53%, the portfolio is ceiling-heavy, and the strictly non-ceiling subset is too small, we use this only as a low-sensitivity frozen-margin surface check rather than a general equivalence result.

Our setting is a source-native replay built from first-party U.S. Bureau of Economic Analysis (BEA), NOAA Climate Prediction Center, and Energy Information Administration (EIA) release histories (U.S. Bureau of Economic Analysis, 2026; NOAA Climate Prediction Center, 2026; U.S. Energy Information Administration, 2026). The ledger contains 35 frozen endpoints and 2,056 retained runtime-valid rows spanning cutoff validity, release alignment, and exogenous grounding. The initial BEA/NOAA object is motivated by TimeSage-EV but is not an exact rerun because the evaluated TimeSage-EV snapshot remains unavailable (Yao et al., 2026). EIA endpoint construction,

the 4/8 temporal split, scorer, procedure hashes, and condition order were frozen before EIA outcomes, but EIA itself was selected post-review because its recurring dated releases are structurally compatible with the cutoff mechanism. We therefore treat it as a post-hoc compatibility probe, not a prospective cross-domain confirmation. Endpoints, not repeated executions, are the independent scientific units.

The paper is an evaluation and identification study, not a skill-learning algorithm. Prior work already studies causal skill attribution, reusable-skill evolution, and temporal retrieval (Zhou et al., 2026; Dong et al., 2026; Wang et al., 2026b; Guan et al., 2026; Wei et al., 2026; Zheng et al., 2026; Xu et al., 2026; Asai et al., 2024; Yan et al., 2024; Han et al., 2025; Wang et al., 2026a). We do not claim novelty in procedure reuse, time-aware retrieval, or paired skill causality. The residual contribution is an audit contract for assigning scientific credit and an empirical demonstration that stronger controls can reverse a repair verdict, expose a surface interaction, downgrade a support cell, and erase container-level credit while leaving some net repairs intact.

Contributions. We make three contributions. First, we separate net repair, same-surface operation-output contrast, placebo/surface perturbation, and integration-surface effect rather than collapsing them into one “skill effect.” Second, we show that this separation changes conclusions: $T > G$ but $T = N$ is an observed false positive; fresh controls preserve bounded DeepSeek repairs while reclassifying EIA as surface-interacting and downgrading Kimi grounding. Third, an exact-output context-materialization control yields a passing TOST at the prospectively frozen ± 10 -point *portfolio-average* margin; because the design has low sensitivity and a ceiling-heavy portfolio, we treat this as a frozen-margin surface check while leaving the low-count non-ceiling subset and broader temporal-retrieval questions unresolved. The contribution is therefore methodological understanding: within the registered endpoint portfolio and tested resolution, stronger controls can preserve, narrow, or reverse an apparent reusable-skill effect.

Reader’s map. Figure 1 motivates the comparator and attribution problem. Figure 2 summarizes the full intervention ladder and shows how the later R17 search-projection study fits as a bounded mechanism child rather than a replacement for the R16 paper. The original causal intervention map and regime-effect plot remain as detailed mechanism/result figures.

2 AN INTERVENTION AUDIT OF REUSABLE PROCEDURES

The audit treats a reusable procedure as an intervention rather than a label attached to a successful run. Let endpoint i specify a question x_i , evidence package E_i , model configuration M_i , and family-specific binary outcome $Y_i \in \{0, 1\}$. For repeated executions under arm a , let $\bar{Y}_i(a)$ be the repeat mean. Repeats measure execution robustness; endpoints remain the independent scientific units.

Base stress test: N, G, and T. The historical design contains the original agent N, a complexity-matched target-blind *stress helper* G, and the targeted procedure T. G is intentionally mechanism-misaligned rather than a benign or optimal generic assistant; its role is to test whether T-versus-structured-assistance can produce a misleading verdict. We report

$$\hat{\Delta}_{T-N} = \frac{1}{|T|} \sum_i (\bar{Y}_i(T) - \bar{Y}_i(N)), \quad \hat{\Delta}_{T-G} = \frac{1}{|T|} \sum_i (\bar{Y}_i(T) - \bar{Y}_i(G)). \quad (2)$$

T-N is the net effect of the targeted intervention relative to the original agent. T-G asks whether it beats the executed target-blind helper. G is not assumed benign, so $T > G$ with $T \leq N$ is no repair.

Same-surface placebo: G_0 . G_0 executes the same helper-exposure path as T but is information-empty: it reads neither package nor context and returns an empty object. We report

$$\hat{\Delta}_{T-G_0} = \frac{1}{|T|} \sum_i (\bar{Y}_i(T) - \bar{Y}_i(G_0)), \quad \hat{\Delta}_{G_0-N} = \frac{1}{|T|} \sum_i (\bar{Y}_i(G_0) - \bar{Y}_i(N)). \quad (3)$$

$T-G_0$ is a *same-surface operation-output contrast*: it asks whether supplying the targeted output changes behavior relative to no operation output under the same exposure surface. G_0-N is the corresponding surface/placebo perturbation diagnostic. These quantities obey the observed-mean identity

$$\hat{\Delta}_{T-N} = \hat{\Delta}_{T-G_0} + \hat{\Delta}_{G_0-N}. \quad (4)$$




Figure 2: **Method: intervention ladder for repair and attribution, with the bounded R17 mechanism child.** The R16 audit keeps four contrasts separate: T-N measures net repair, T-G₀ is the same-surface operation-output contrast, G₀-N diagnoses helper-surface perturbation, and T-R_{surf} tests integration surface while holding the exact operation output fixed. Fail-closed rules prevent T>G with T≤N from being called repair and prevent nonzero G₀-N from being hidden inside a “pure operation” claim. The lower branch summarizes R17: WIN-C and MRW expose different views of the same realized search pool to a persistent updater. Its global effect is unresolved, one frozen First-Fail state is stable, and independently generated updater states do not reproduce that benefit; R17 therefore remains bounded mechanism evidence rather than a standalone promotion claim.

Equation 4 is a decomposition of contrasts, not a mediation theorem. A minimal causal bookkeeping view has an exposure surface S , operation output O , and outcome Y : N leaves the experimental helper absent, G₀ exposes the helper surface with $O = \emptyset$, T exposes the same surface with the frozen targeted output, and R_{surf} holds that output fixed while changing where it enters the harness. This intervention map clarifies which variable each contrast changes, but it is not a full SCM identification result and makes no mediation claim. When G₀-N is nonzero, T-G₀ is conditional on a perturbed surface and must not be called a pure operation effect. Accordingly, we classify T-N > 0 as *net repair*; T-G₀ is reported separately as evidence about the output under that surface. Cells with nonzero placebo shift are explicitly labeled surface-interacting rather than operation-attributed.

Integration-surface control: R_{surf}. R_{surf} runs the exact same frozen operation source on the exact same candidate evidence as T, then materializes the identical operation output inside ordinary context while exposing no callable helper. Thus

$$\hat{\Delta}_{T-R_{\text{surf}}} = \frac{1}{|\mathcal{I}|} \sum_i (\bar{Y}_i(T) - \bar{Y}_i(R_{\text{surf}})) \quad (5)$$

is an integration-surface estimand under the frozen forced one-answer harness. R_{surf} is deliberately not an independently learned or heuristic temporal retriever. A zero-call parity audit verifies that candidate evidence is preserved, operation-output content is identical, and only where the pre-computed output enters the harness changes. The prospectively frozen material threshold is 10 percentage points on the pre-frozen DeepSeek endpoint portfolio. Multi-turn deployments in which a callable helper can be invoked adaptively require a different control.

The audit targets three recurring failure families. **Cutoff validity** filters evidence by the target information cutoff. **Release alignment** maps a release to the reporting period represented by its content. **Exogenous grounding** links a numerical change to first-party documentary evidence available before the cutoff. No targeted procedure contains endpoint-specific answers, gold labels, or fitted values; procedure hashes are frozen before execution and reused on held-out releases.

Three falsifiers follow. First, T>G with T≤N is no repair. Second, a positive T-G₀ does not by itself establish a pure operation effect when G₀-N is nonzero; all three contrasts in Eq. 4 must remain




Figure 3: Base three-arm stress test. N anchors the original agent, G is a target-blind mechanism-misaligned stress helper, and T contains the targeted temporal operation. G_0 exposes placebo/surface perturbation; R_{surf} isolates where the same precomputed operation output enters the forced one-answer harness.

visible. Third, if the integration surface matters under the forced one-answer harness, T should leave a material residual over R_{surf} when operation output and evidence are held exactly fixed.

3 RELATED WORK

Causal skill attribution and agent evaluation. Paired trace auditing, no-skill comparisons, randomized skill masking, and continual skill evaluation already show that exposing an agent to a skill can change behavior and that effects are heterogeneous (Zhou et al., 2026; Dong et al., 2026; Wang et al., 2026b; Guan et al., 2026). We therefore do not claim paired skill causality or heterogeneous effects as new. Our question is narrower and evaluative: does the targeted intervention improve the original agent, does the target operation output matter under the same exposure surface, and does moving the same precomputed output between two integration surfaces change the final answer under a forced one-answer harness? The observed $T > G$ but $T = N$ cutoff cell demonstrates why these are different estimands rather than interchangeable baselines.

Skill evolution and reusable procedures. Evo-Harness, SkillProx, HyperSkill, Anything2Skill, and related systems study continual maintenance, compilation, optimization, structured skill memory, and conversion of external knowledge into reusable procedures (Wei et al., 2026; Zheng et al., 2026; Xu et al., 2026; Pan et al., 2026). These works make procedure reuse and skill accumulation established objects. We treat acquisition as upstream: after a reusable procedure exists, our contribution is an audit of the scientific credit assigned to it. A stronger control may preserve a net repair while exposing a surface interaction, eliminating a skill-container interpretation, or downgrading a previously positive support cell.

Temporal retrieval and context construction. Selective/corrective RAG and temporal retrieval already address irrelevant evidence, time-sensitive indexing, and explicit temporal constraints upstream of generation (Asai et al., 2024; Yan et al., 2024; Zhang et al., 2025; Abdallah et al., 2026; Han et al., 2025; Wang et al., 2026a; Liu et al., 2026). Consequently, a useful cutoff or alignment operation need not be packaged as callable state. Our R_{surf} experiment does *not* compare T with an independently learned temporal retriever. It intentionally runs the same frozen operation and materializes the identical output in context, creating a same-information surface-isolation test. The result therefore speaks only to integration-surface placement under the frozen one-answer harness, not to algorithmic equivalence between skills and temporal RAG or to adaptive callable use in deployment. A broader retrieval claim would require a separately implemented retriever or reranker that must discover the relevant transformation from the same raw candidates; we do not make that claim here.

The residual novelty is thus methodological: a reusable-skill evaluation should distinguish net repair, same-surface output contrast, placebo/surface perturbation, and integration-surface value before assigning credit to “the skill.” In our data, making those distinctions changes multiple verdicts rather than merely adding confidence intervals to an unchanged story.

4 INTERVENTIONS AND AUDIT CONTRACTS

All arms replay the same frozen endpoint with the same evidence package, model configuration, output schema, and family scorer. We intervene only on the procedure exposure or on where an already computed procedure output enters the harness. The scientific unit is the endpoint; repeated executions are robustness observations rather than additional tasks.

4.1 BASE N/G/T STRESS TEST

N adds no experimental helper. G exposes a frozen target-blind *stress helper* using the same callable surface and matched implementation complexity as T, but it is forbidden from performing the registered temporal operation. G is not presented as a benign, optimal, or representative generic assistant: its purpose is to stress the common T-versus-generic evaluation pattern and reveal when control-arm degradation changes the verdict. For cutoff validity, G can validate package structure but cannot compare release dates with the cutoff; for release alignment, it cannot map publication dates to reporting periods; for grounding, it cannot perform the mechanism-specific contextual selector. Targeted and G source implementations are matched in lexical-token and AST-node count for each family.

This base test is useful precisely because G need not be benign. T-G is a stress contrast against structured but mechanism-misaligned assistance, while T-N determines whether anything was actually repaired. We do not use T-G alone to estimate a generic-helper treatment effect. The DeepSeek cutoff false positive in Section 6 shows why the N anchor is mandatory whenever such a stress comparator is used.

4.2 SAME-SURFACE NO-OPERATION PLACEBO G_0

G_0 executes the same assigned-helper path as T but its frozen function is simply `def skill(package, context): return {}`. It reads neither evidence nor task context. Consequently T and G_0 differ in whether the target operation output is present while sharing the same helper-exposure surface. We do not label G_0 behaviorally neutral at the stratum level. Its empirical difference from N is reported directly.

The final audit therefore does not combine T-N and T- G_0 into a single pure-operation estimand. T-N carries the *net-repair* claim. T- G_0 is a *same-surface operation-output contrast*; G_0 -N reports how much the placebo/surface itself perturbs N. Their observed means satisfy Eq. 4. When G_0 harms or helps N, the T- G_0 contrast is conditional on that perturbed surface and is reported as such. The separately frozen A1/A2 G_0 studies remain diagnostics of wrapper perturbation; their full-track equivalence result is not transported into every small load-bearing stratum.

4.3 CONTEXT-MATERIALIZATION SURFACE CONTROL R_{surf}

R_{surf} isolates a narrower implementation question. It runs the exact targeted-operation source on the same candidates and computes the exact same operation output as T. Instead of exposing that output as `reusable_helper_output`, it adds the object as `retrieval_operation_output` inside the ordinary evidence context and exposes no callable helper. A zero-call audit verifies for every tested endpoint that removing this one added field reconstructs the original evidence package, operation-output hashes match T exactly, and candidate evidence is neither removed nor reordered.

Because R_{surf} deliberately reuses T's operation output, it is a *surface-isolation control*, not an independent temporal-retrieval algorithm. It asks whether placing the same precomputed output on the T surface rather than in ordinary context changes the final answer under this forced one-answer harness. It does not test adaptive callable use during generation, whether a learned temporal retriever can discover or approximate the operation, or whether persistent callable state has an advantage in multi-turn deployment.

4.4 FROZEN OUTCOMES, THRESHOLDS, AND EVIDENCE LAYERS

Before each source-native execution tranche, endpoint identity, family outcome, repeat count, condition order, model configuration, and procedure source hashes are frozen. A cutoff success requires a

nonempty cited set and latest selected reference that are all available by the target cutoff. A release-alignment success requires the correct reporting-period identity and aligned value. A grounding success requires the requested evidence slots to be filled by frozen first-party sentence IDs in the correct order. Auxiliary stricter scorers are reported only as sensitivity analyses.

The original R_{surf} contract freezes a 10-point material-effect threshold on the 18-endpoint DeepSeek portfolio before any R_{surf} outcome is observed. The contract records this as a prospective *materiality* threshold, not a power-derived equivalence margin; it does not contain a prospective variance-based power justification. We therefore do not retrofit one. As post-review context only, the targeted contrasts already available before the R_{surf} run were on a 20–60 point scale, making 10 points a conservative practical scale for a residual worth detecting, but this scale argument is not presented as an additional preregistered fact. We report the preregistered 90% endpoint-bootstrap interval, a post-hoc 95% bootstrap sensitivity interval, and a paired-endpoint TOST at the unchanged ± 10 -point margin. A separate post-review power calculation is used only to plan future evidence (Appendix A.8), not to upgrade the completed experiment. For G_0 , A1 and A2 remain unpooled; exact-stratum G_0 –N diagnostics show where full-track equivalence does and does not transport.

The initial BEA/NOAA replay freezes pre-April and April–May periods before outcomes. A separate post-hoc EIA compatibility probe uses 16 fixed weekly releases to construct 12 deterministic cutoff endpoints: source indices 3–14, with four releases before April 1 assigned to C3–R4 and eight later releases assigned to C4–R4. Endpoint construction, scorer, T1/G1 hashes, and randomized condition order were frozen before EIA model outcomes. However, EIA was selected post-review because its dated weekly releases are structurally compatible with the cutoff mechanism, so C4–R4 is an illustrative mechanism-compatibility probe rather than prospective cross-domain confirmation.

5 EXPERIMENTAL SETUP: SOURCE-NATIVE REPLAY

TimeSage-EV motivates the three recurring mechanism families—cutoff validity, release alignment, and exogenous grounding—and exposes persistent skills to later periods (Yao et al., 2026). Our experiment is not an exact rerun of its unavailable evaluated snapshot. Instead, we reconstruct content-addressed endpoints directly from first-party institutional release histories and intervene on the same reusable operations.

5.1 FROZEN ENDPOINTS AND SCORERS

The initial replay contains 23 endpoints from BEA GDP releases and NOAA-CPC ENSO Diagnostic Discussions (U.S. Bureau of Economic Analysis, 2026; NOAA Climate Prediction Center, 2026): twelve pre-April endpoints form C3–R and eleven April–May endpoints form held-out C4–R. Raw pages are SHA-256 bound, evidence identifiers are opaque, and grounding exposes the full content-region sentence pool rather than a gold-only candidate set. Family scorers are deterministic and frozen before model outcomes; a pre-execution self-test gives 23/23 success on gold outputs and 0/23 on constructed negatives.

A post-review compatibility probe adds 12 EIA Weekly Petroleum Status Report cutoff endpoints from a third institutional system (U.S. Energy Information Administration, 2026). EIA is chosen because its recurring first-party releases and explicit publication dates fit the already frozen cutoff interface. Four endpoints form an earlier diagnostic split and eight April–May endpoints form the later probe split. T1 and G1 are reused byte-for-byte; source manifest, endpoints, scorer, harness, and condition order are frozen before execution. Because the institution itself was selected post-review for mechanism compatibility, EIA is an illustrative compatibility probe rather than independently preregistered cross-domain confirmation.

5.2 QUESTIONS AND INFERENCE CONTRACT

RQ1 asks whether adding the original-agent anchor changes a verdict that would look positive under targeted-versus-generic comparison alone. RQ2 asks whether the targeted operation output improves on a same-surface no-operation placebo G_0 while the complete targeted intervention also improves on N. G_0 –N is reported separately as a surface-perturbation diagnostic rather than assumed neutral within every stratum. RQ3 asks whether byte-identical procedures retain direction on compatible

held-out releases while preserving ceiling and incompatible-domain nulls. RQ4 asks whether moving the exact same frozen temporal-operation output between the T surface and ordinary context leaves a material residual under the forced one-answer harness.

The endpoint is the independent scientific unit. Repeats characterize execution robustness and are averaged within endpoint; they are never counted as additional tasks. For non-ceiling cells we report endpoint win/tie/loss patterns, percentile bootstrap intervals over endpoints, and exact two-sided sign tests over non-ties when resolution permits. Two- and three-endpoint cells are explicitly descriptive. We do not pool heterogeneous procedures, models, or domains into an omnibus procedure-effect statistic.

5.3 EVIDENCE LAYERS AND PARITY

DeepSeek-V4-Pro is the registered primary track with five-repeat historical replay plus fresh identification runs. A pre-outcome support rule selects Kimi-K3 as a separately labeled replication; a later hash audit removes one overwritten historical repeat block without replacement. Doubao-Seed-2.0-Mini is an exploratory capacity stress. The EIA layer evaluates DeepSeek and the mini model. After adding the prospectively frozen G_0 and R controls, the retained source-native evidence ledger contains 2,056 runtime-valid rows over the same 35 distinct endpoints from BEA, NOAA, and EIA. The endpoint count, not the enlarged repeat/control-row count, determines scientific resolution.

Every targeted procedure implementation has a family-matched generic helper with the same callable signature, identical Python lexical-token count and AST-node count, and equal external tool-call budget. G_0 uses the same helper-execution surface but is information-empty: it returns $\{\}$ and reads neither package nor context. For R_{surf} , zero-call parity verifies on every tested endpoint that removing the added context field recovers the original evidence package exactly and that the materialized operation output is content-identical to T’s helper output. This makes R_{surf} a controlled integration-surface intervention, not an independently designed retriever. Full execution chronology, content hashes, per-call CSV checkpoints, and provenance bindings are reported in Appendix A and the supplement.

6 RESULTS

We report endpoint-level contrasts and keep repeated executions inside each endpoint mean. The results are organized by what each added control changes about the scientific verdict rather than by a universal “skills help” narrative.

6.1 THE ORIGINAL-AGENT ANCHOR OVERTURNS A TWO-ARM GAIN

On DeepSeek C3-R cutoff, T succeeds on 100% of endpoint-repeat observations and the complexity-matched target-blind stress helper G on 72%. A conventional T-versus-G comparison therefore reports +28 points. N is also 100%, however, so $T-N=0$ and the cell is *no repair*. At endpoint means, T-G has 4 wins/1 tie/0 losses, whereas T-N has 0 wins/5 ties/0 losses, so the verdict reversal does not depend on treating repeats as independent tasks. A zero-call audit further confirms equal helper invocation and no detectable condition-position association; T prompts are 18.85 input tokens shorter on average. The point is not that G is a universally fair generic assistant, but that T-versus-stress-helper evidence can look convincingly positive while the original agent was never repaired.

6.2 G_0 IS A PLACEBO DIAGNOSTIC, NOT A UNIVERSAL NEUTRALITY ASSUMPTION

The independently frozen A2 study shows that G_0 is close to N when averaged over the full DeepSeek endpoint portfolio ($G_0-N=-2.1$ points, 90% CI $[-9.3, 5.0]$); Kimi is similarly close on its full endpoint set (+2.2 points, $[-4.3, 8.7]$). These global summaries do *not* transport into every small load-bearing stratum. On the exact A2 C3 and C4 grounding strata, G_0-N is +5 points and the 90% intervals extend to +15 points; on the eight EIA held-out endpoints it is -18.8 points with a wide interval. We therefore remove stratum-level “behaviorally neutral” language and use G_0 as a same-surface no-operation placebo.




Figure 4: Temporal-operation effects are regime-dependent. Circles compare T with N; crosses compare T with historical target-blind G. DeepSeek C3 cutoff illustrates the identification point: $T > G$ but $T = N$, so the apparent two-arm gain is not a repair. EIA is a post-hoc compatibility probe on a separate institutional endpoint set.

The three contrasts are more informative when kept separate. In a fresh two-repeat tranche, DeepSeek C3 grounding has $T-N=+30$, $T-G_0=+30$, and $G_0-N=0$ points (2 endpoint wins, 3 ties, 0 losses for the first two contrasts); C4 grounding gives $+20$, $+20$, and 0 (1 win, 4 ties, 0 losses). The endpoint-bootstrap intervals include zero because each cell has five endpoints. These are directional net-repair and same-surface output results with zero *observed* placebo shift, not formal stratum-level equivalence or pure-operation identification.

The EIA cell provides a compact running example of the whole audit. On eight post-hoc compatibility-probe endpoints, fresh $T-N$ is $+37.5$ points with 90% bootstrap interval $[18.8, 56.3]$, five wins, three ties, no losses, and one-sided exact sign test $p = 0.03125$. $T-G_0$ is $+62.5$ points while G_0-N is -25 points, giving the exact contrast decomposition $+37.5 = +62.5 - 25$. In the later exact-output surface check, T and R_{surf} are both 100% on these same eight endpoints. Thus this single example reads as: net repair over N, a larger same-surface contrast partly created by a harmful placebo shift, and no observed integration-surface residual at ceiling. We classify it as *net repair with surface interaction*, not confirmatory cross-domain evidence or pure-operation attribution. The earlier five-repeat EIA execution is directionally consistent, but both executions reuse the same endpoint set and are not independent endpoint samples.

6.3 GROUNDING SURVIVES ONLY AS A BOUNDED DEEPSEEK RESULT

The historical DeepSeek grounding cells are 20/20/60 for N/G/T in C3 and 40/40/60 in C4. Fresh N/ G_0 /T execution preserves $+30$ and $+20$ point net-repair and same-surface output directions with zero observed G_0-N shift in those tranches. Kimi does not supply cross-model grounding confirmation: its fresh same-domain grounding check has $T-G_0 = -25$ points and $T-N=0$, and cross-domain grounding is also null. We therefore drop the earlier cross-model grounding-transfer interpretation and retain only a bounded DeepSeek grounding result.

6.4 EIA IS A POST-HOC COMPATIBILITY PROBE

The EIA expansion is constructed from 16 fixed weekly releases. Source indices 3–14 deterministically define 12 endpoints; releases before April 1 form four earlier diagnostics and the remaining eight form C4-R4. The endpoint set, time split, scorer, T1/G1 hashes, and randomized condition order are frozen before EIA outcomes. But EIA itself was chosen post-review because its dated releases fit the already frozen cutoff interface. We therefore use it only as a mechanism-compatibility

illustration, not as confirmatory or generic cross-domain evidence. The mini model is already at N ceiling on the same eight endpoints, further emphasizing regime dependence.

6.5 EXACT-OUTPUT CONTROL PASSES A LOW-SENSITIVITY FROZEN-MARGIN CHECK

R_{surf} deliberately computes the exact T operation output and changes only how that output enters the harness. Across 18 pre-frozen DeepSeek endpoints, $T-R_{\text{surf}}$ is 0.0 points (1 win/16 ties/1 loss). The preregistered 90% bootstrap interval is $[-5.6, 5.6]$; a paired TOST at the prospectively frozen materiality margin of ± 10 points passes (both one-sided $p = 0.0121$), with paired-t 95% CI $[-8.5, 8.5]$ and bootstrap 95% sensitivity CI $[-8.3, 8.3]$. The margin was fixed before R outcomes and was not derived from the observed $T-R_{\text{surf}}$ variance. We treat the result as a *low-sensitivity frozen-margin check*, not positive evidence that the surfaces are generally interchangeable: an observed-variance calculation gives only about 53% probability of declaring equivalence at $n = 18$ if the true mean were zero, and the portfolio is ceiling-heavy.

The zero mean is not produced only by ceiling cells: C3 grounding is 60%/60% under T/R_{surf} , and C4 is 70%/70%. However, the four strictly non-ceiling endpoints have a wide 90% interval $[-25, 25]$, so non-ceiling equivalence remains unresolved; EIA is ceiling at 100%/100%. All 18 endpoint residuals are listed in Appendix 5. Only two endpoint means are non-ties, both in C4 grounding, and they have opposite signs: one is +50 points because only T repeat 0 selects the registered successful span pair, while the other is -50 points because only R_{surf} repeat 0 does so. In both endpoints the frozen operation-output object is content-identical across T and R_{surf} ; the other three arm-repeat outcomes at that endpoint follow the failing selection. This pattern is consistent with low-frequency generation/selection variability under different placement surfaces, not a stable one-direction surface advantage. A three-endpoint Kimi alignment tie is descriptive only. Because R_{surf} reuses T's exact output under a forced one-answer harness, no multi-turn callable-state or independent temporal-retrieval equivalence is implied.

6.6 BOUNDED MECHANISM UPDATE: SERVING PROJECTION DOES NOT REPRODUCIBLY GENERATE THE BENEFICIAL UPDATER STATE

A later controlled child study asks a narrower learning-interface question: holding the realized candidate search pool fixed, does the evidence view supplied to a persistent updater matter? We compare a winner-only view (WIN-C) with a richer memory-relevant view (MRW) over 48 paired units arranged as 12 streams with four replicates per stream. The paired mean difference is $\text{MRW} - \text{WIN-C} = +0.023148$. Under the frozen analysis contract, neither superiority nor practical equivalence is established, so the global projection effect remains unresolved.

The child study nevertheless separates state stability from treatment reproducibility. One exact historical First-Fail updater state remains beneficial when evaluated again: fresh evaluation replicates change the held-out outcome by $+1/18$ and $+4/18$. However, when the same rendered First-Fail evidence is passed through independently generated updater states, the corresponding differences are $0/18$ and $-1/18$. The replication gate therefore fails. The observed positive state is a real frozen-state existence/stability witness, but First-Fail evidence did not reproducibly generate that state under the tested updater.

This result does not alter the R16 repair/attribution verdicts above and does not justify rescue by adding providers, models, or seeds. Its bounded mechanism lesson is that evidence serving and updater-state generation are distinct control surfaces: even if a particular learned state is stable once obtained, a serving projection deserves causal credit only when it reproducibly changes the distribution of generated updater states. In the present R17 object, updater-state generation is the identified precision/reproducibility bottleneck and standalone promotion remains unsupported.

7 LIMITATIONS

The frozen TimeSage-EV evaluated snapshot remains unavailable, so the 35-endpoint BEA/NOAA/EIA study is independent source-native evidence rather than an exact rerun. Resolution is uneven: the main grounding cells have five endpoints, Kimi alignment three, Kimi same-domain grounding two, and cross-domain grounding is null. G_0 is an information-empty same-surface

placebo, not a universally neutral control: full-track equivalence does not transport to every stratum, and EIA G_0 underperforms N. Accordingly, T-N carries net repair while T- G_0 and G_0 -N remain separate output/surface contrasts.

R_{surf} tests only placement of the exact same precomputed output under a forced one-answer harness. The portfolio TOST passes at the frozen ± 10 -point margin, but sensitivity is about 53%, the portfolio is ceiling-heavy, and four strictly non-ceiling endpoints remain unresolved; we therefore treat it as a low-sensitivity margin check, not positive evidence of general surface equivalence. Adaptive callable use and independent temporal retrievers are untested. EIA is a post-hoc compatibility probe: its endpoint construction and temporal split were outcome-independent, but the institution was not independently sampled, so it does not support prospective cross-domain confirmation or model-invariant transfer. The later R17 serving-projection child is also deliberately bounded: its 48 paired units do not establish MRW superiority or practical equivalence, and the beneficial First-Fail state is stable under reevaluation but is not reproducibly regenerated from the same evidence. We therefore treat R17 as mechanism/negative evidence about updaters-state generation, not as an additional positive result or a reason to reopen the completed R16 experiment.

8 CONCLUSION

A reusable “skill effect” can contain several different causal stories. In our audit, the original-agent anchor turns one apparent +28-point targeted-versus-stress-helper benefit into no repair; a same-surface no-operation placebo exposes output and surface effects without requiring neutrality; and an exact-output context-materialization control makes the observed portfolio consistent with no integration-surface difference larger than the frozen ± 10 -point margin under the forced one-answer harness. Stronger controls also force a Kimi grounding downgrade while leaving only bounded, low-resolution positive repair evidence.

The resulting contribution is therefore not that temporal skills are generally superior. It is an intervention audit for assigning credit: use T-N for net repair; keep T- G_0 and G_0 -N separate to expose same-surface output and placebo/surface effects; and use T- R_{surf} only as an integration-surface test when the same precomputed output is placed on two surfaces under the same one-answer harness. In this study, stronger controls preserve some bounded net repairs, expose a surface interaction, reverse support cells, and produce a passing but low-sensitivity frozen-margin surface check. Broader claims about non-ceiling equivalence, adaptive callable interfaces, or temporal retrieval remain outside the evidence.

REPRODUCIBILITY STATEMENT

The source-native study is content-addressed end to end. The supplementary replication bundle contains the frozen endpoint definitions, prompt harnesses, family scorers, all six executed targeted/generic helper implementations, condition schedules, first-party archived source pages, retained raw result JSON, adjudication code, and a SHA-256 manifest. Appendix A records the executed helper code, the endpoint-level inferential protocol, and the zero-call invocation/order/date-serialization audits. The source-native bundle is independent of the unavailable TimeSage-EV evaluated snapshot; the exact TimeSage replication remains a separately identified external replication target.

AI USE STATEMENT

Generative AI tools were used throughout the research workflow to assist literature triage, research-idea and hypothesis refinement, experimental-design critique, implementation and debugging, data reformatting, result interpretation, internal mock review, and manuscript drafting and editing. AI-assisted experimental artifacts were accepted only after executable checks against frozen task contrasts, immutable raw outputs, first-party source material, or deterministic recomputation. AI-generated review feedback had advisory status and did not itself create scientific evidence. The authors reviewed the AI-assisted work and take responsibility for the final paper, claims, code, and supplementary artifacts.

REFERENCES

- Abdelrahman Abdallah, Bhawna Piryani, Jonas Wallat, Avishek Anand, and Adam Jatowt. Tempretriever: Fusion-based temporal dense passage retrieval for time-sensitive questions. In *Proceedings of the 19th ACM International Conference on Web Search and Data Mining*, pp. 5–15. Association for Computing Machinery, 2026. doi: 10.1145/3773966.3777938.
- Akari Asai, Zeqiu Wu, Yizhong Wang, et al. Self-rag: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*, 2024.
- Gen Dong, Yanjie Gao, Liquan Li, et al. Agent skills can be harmful: An empirical study of skill-induced failures in llm agents, 2026.
- Tianyi Guan, Yiding Wang, Haotong Yang, et al. Continualskillbench: Can llm agents truly evolve their capabilities?, 2026.
- Jiale Han, Austin Cheung, Yubai Wei, et al. Rag meets temporal graphs: Time-sensitive modeling and retrieval for evolving knowledge, 2025.
- Hanbing Liu, Lang Cao, and Yang Li. RAG or learning? understanding the limits of LLM adaptation under continuous knowledge drift in the real world. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 11234–11252. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.findings-acl.546. URL <https://aclanthology.org/2026.findings-acl.546/>.
- NOAA Climate Prediction Center. ENSO Diagnostic Discussion Archive. https://www.cpc.ncep.noaa.gov/products/expert_assessment/ENSO_DD_archive.php, 2026. First-party monthly discussions accessed 2026-08-22.
- Qianjun Pan, Yutao Yang, Junsong Li, Jie Zhou, Kai Chen, Xin Li, Qin Chen, and Liang He. Anything2skill: Compiling external knowledge into reusable skills for agents, 2026.
- U.S. Bureau of Economic Analysis. GDP News Release Archive. <https://www.bea.gov/news/archive>, 2026. First-party historical releases accessed 2026-08-22.
- U.S. Energy Information Administration. Weekly Petroleum Status Report Archive. <https://www.eia.gov/petroleum/supply/weekly/archive/>, 2026. First-party weekly release archive accessed 2026-08-22.
- Xiaoman Wang, Yaoze Zhang, Wenzhuo Fan, et al. Ia-rag: Interval-algebra-driven temporal reasoning for dynamic knowledge retrieval, 2026a.
- Yixuan Wang, Yiyang Zhou, Yiming Liang, et al. Not all skills help: Measuring and repairing agent knowledge, 2026b.
- Tianxin Wei, Zhan Shi, Minhua Lin, et al. Evo-harness: Context-to-harness skill compilation for self-evolving agents, 2026.
- Ruiyao Xu, Tiankai Yang, and Wei-Chieh Huang. Hyperskill: Self-evolving llm agents via hypergraph-structured skill memory, 2026.
- Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective retrieval augmented generation, 2024.
- Qingren Yao, Yaxuan Kong, Yuqi Nie, Yichen Li, Anna Vettoruzzo, Qingsong Wen, Ming Jin, Joaquin Vanschoren, and Stefan Zohren. Timesage-ev: A live benchmark for agentic time series analysis in evolving environments, 2026.
- Siyue Zhang, Yuxiang Xue, Yiming Zhang, Xiaobao Wu, Anh Tuan Luu, and Chen Zhao. MRAG: A modular retrieval framework for time-sensitive question answering. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 3080–3118. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.findings-emnlp.167. URL <https://aclanthology.org/2025.findings-emnlp.167/>.

Mingxuan Zheng, Yujin Zhou, Chuxue Cao, et al. Skillprox: Self-evolving agent skills via proximal textual gradient descent, 2026.

Xiaolin Zhou, Jinbo Liu, Li Li, et al. Counterfactual trace auditing of llm agent skills, 2026.

A REPRODUCIBILITY AND SENSITIVITY AUDITS

T1/G1 are byte-identical across the initial replay and the post-hoc EIA compatibility probe. The provenance-clean r3 reaudit is 6221016166b71ebe; r4 adjudication is 316ee89a8eb1d53f. The r5 replication bundle contains executable r3/r4 endpoints, harnesses, scorers, runner scripts, frozen helper source, first-party archived pages, retained result JSON, and a SHA-256 manifest. This bundle reproduces the source-native study and remains separate from the unavailable TimeSage-EV evaluated snapshot.

A.1 EXACT FROZEN REUSABLE HELPERS

The six helper implementations below are the executed source, not pseudocode. Each targeted/generic pair has equal lexical-token count and equal AST-node count under the pre-execution audit.

```
# T1_temporal_cutoff.py
def skill(package, context):
    limit = context.get("cutoff_date", "")
    chosen = []
    for item in package.get("evidence_metadata", []):
        marker = item.get("release_date", "")
        ref = item.get("evidence_ref", "")
        if isinstance(marker, str) and marker <= limit:
            chosen.append(ref)
    return {"evidence_refs": chosen}

# G1_temporal_cutoff.py
def skill(package, context):
    limit = context.get("schema_ceiling", "")
    chosen = []
    for item in package.get("evidence_metadata", []):
        marker = item.get("schema_version", "")
        ref = item.get("evidence_ref", "")
        if isinstance(marker, str) and marker <= limit:
            chosen.append(ref)
    return {"evidence_refs": chosen}

# T2_release_alignment.py
import re
def skill(package, context):
    text = " ".join(package.get("document_text", "").split())
    match = re.search(r"([1-4])?(?:st|nd|rd|th) Quarter(?: and Year)? (20\d{2})",
        text, re.I)
    value = None
    if match:
        value = f"{match.group(2)}-Q{match.group(1)}"
    return {"value": value}

# G2_release_alignment.py
import re
def skill(package, context):
    text = " ".join(package.get("document_text", "").split())
    match = re.search(r"(News) (Release|Archive|Statement)(?: and Data)? "
        r"([A-Za-z]{2,20})", text, re.I)
    value = None
    if match:
        value = f"{match.group(1)}-{match.group(3)}"
    return {"value": value}

# T3_exogenous_grounding.py
def skill(package, context):
    markers = ("contributors to", "reflected", "because", "collectively",
        "favored", "favors ", "associated with", "due to")
    chosen = []
    for span in package.get("candidate_spans", []):
        text = span.get("text", "").lower()
        sid = span.get("span_id", "")
        if any(marker in text for marker in markers):
            chosen.append(sid)
```

```

702     return {"span_ids": chosen}
703
704 # G3_exogenous_grounding.py
705 def skill(package, context):
706     markers = ("table", "figure", "release", "appendix", "contact",
707               "notes", "technical notes", "page")
708     chosen = []
709     for span in package.get("candidate_spans", []):
710         text = span.get("text", "").lower()
711         sid = span.get("span_id", "")
712         if any(marker in text for marker in markers):
713             chosen.append(sid)
714     return {"span_ids": chosen}
715

```

A.2 EXPLORATORY CAPACITY-STRESS MATRIX

Split	Family	N0	Generic	Target	Δ_{T-G}	Δ_{T-N}
C3-R	Cutoff	32	0	72	+72	+40
C3-R	Align	100	100	100	0	0
C3-R	Ground	40	48	64	+16	+24
C4-R	Cutoff	53	13	100	+87	+47
C4-R	Align	80	93	100	+7	+20
C4-R	Ground	60	52	64	+12	+4

Table 1: Exploratory mini-model capacity stress over five repeats (percent), reported separately from the registered primary and replication tracks.

A.3 ENDPOINT-LEVEL UNCERTAINTY AND TEST RESOLUTION

Table 2 averages repeats within endpoint-condition and bootstraps endpoints; W/T/L is targeted versus generic, and sign tests use non-ties. Repeats never increase the inferential sample size. For n non-tied endpoints, the smallest attainable two-sided sign-test value is 2^{1-n} when every endpoint moves in the same direction. Thus two non-tied endpoints cannot yield a value below .5; three cannot go below .25; five cannot go below .0625; eight can reach .0078. Ties reduce resolution further. This arithmetic is why the two-endpoint grounding-transfer result is explicitly descriptive and why the three-endpoint Kimi alignment result remains underpowered despite a large effect. The eight-endpoint EIA probe has finer endpoint-level resolution than those tiny cells, but because the institution was selected post-review for mechanism compatibility it remains an illustrative compatibility probe rather than confirmatory cross-domain evidence.

Layer	Split	Fam.	Δ_{T-G} [95%]	W/T/L	p_G	Δ_{T-N} [95%]	p_N
D	C3	Cut	28 [8,56]	4/1/0	.125	0 [0,0]	1
D	C3	Grd	40 [0,80]	2/3/0	.50	40 [0,80]	.50
D	C4	Cut	13 [0,20]	2/1/0	.50	0 [0,0]	1
D	C4	Grd	20 [0,60]	1/4/0	1	20 [0,60]	1
K	C3	Grd	10 [0,20]	2/3/0	.50	10 [0,30]	1
K	C4	Align	66.7 [50,75]	3/0/0	.25	41.7 [25,75]	.25
K	C4	Grd	20 [0,60]	1/4/0	1	20 [0,60]	1
M	C3	Cut	72 [32,100]	4/1/0	.125	40 [8,72]	.25
M	C3	Grd	16 [-12,60]	1/3/1	1	24 [0,64]	.50
M	C4	Cut	86.7 [60,100]	3/0/0	.25	46.7 [20,100]	.25
M	C4	Align	6.7 [0,20]	1/2/0	1	20 [0,60]	1
M	C4	Grd	12 [0,28]	2/3/0	.50	4 [-32,36]	1

Table 2: Endpoint-level uncertainty for non-ceiling initial-replay cells. D=DeepSeek primary, K=Kimi replication, M=exploratory mini.

A.4 WHY THE FAMILY ESTIMAND IS PRIMARY

The family contract and the full-task conjunctive score answer different questions. The family score asks whether the targeted temporal operation was corrected. The conjunctive score additionally requires orthogonal output fields and exact surface forms. A zero conjunctive score can therefore coexist with a successful cutoff intervention when temporal admissibility is correct but an unrelated field differs, or with successful release alignment when a semantically identical period is formatted

differently. The family definitions were frozen before execution, so selecting them is determined by the experimental question and its time order. The conjunctive score is preserved as a sensitivity view and is never used to erase an unfavorable family result. Empirically, the trade-off is sharp: grounding is unchanged because its registered evidence slots coincide with the conjunctive requirement, whereas cutoff and alignment collapse to zero under the full-task exact score because that score additionally demands orthogonal fields and exact surface forms. This collapse therefore diagnoses an estimand change rather than showing that the registered temporal operation failed under a stricter version of the same scorer.

Family	r3 full-task conjunctive score	Effect on claim
Grounding	Identical to family score	Unchanged
Cutoff	All cells collapse to 0	Family estimand remains primary
Alignment	All cells collapse to 0	Family estimand remains primary

Table 3: Sensitivity to the auxiliary full-task exact score. The two columns evaluate different estimands.

A.5 CALL-RATE, PROMPT-LENGTH, AND CONDITION-POSITION AUDIT

The retained source-native rows contain 442 no-skill, 442 generic, and 442 targeted observations. Helper execution is deterministic in the frozen harness: generic and targeted execute exactly once before the model answer in 442/442 rows each, whereas no-skill executes no helper in 0/442 rows. Thus targeted-minus-generic cannot be explained by a higher helper invocation rate; this experiment does not measure endogenous tool-use propensity. Across exact targeted/generic pairs, targeted prompts are shorter by 18.85 input tokens on average (302 shorter, 140 longer). Condition-position counts are 157/135/150 for no-skill, 145/155/142 for generic, and 140/152/150 for targeted across positions 0/1/2. Pooled 3×2 success-by-position tests give $p = .993, .657$, and $.914$ for no-skill, generic, and targeted respectively; targeted-minus-no-skill remains positive at all three positions. These are post-hoc zero-call diagnostics and do not substitute for a prospectively counterbalanced replication. The full audit is included in the supplement.

A.6 ISO DATE AND PERIOD-REPRESENTATION AUDIT

A zero-call static audit covers the 23 initial BEA/NOAA endpoints and the 12 post-review EIA endpoints. Across the serialized endpoint objects, all 55 occurrences of `cutoff_date` and all 176 occurrences of `release_date` are calendar-valid YYYY-MM-DD strings. Whenever both top-level and `skill_context` cutoff fields are present they agree exactly. T1 makes 88 observed release-versus-cutoff comparisons over the frozen endpoint packages; lexical string order and parsed-calendar order agree on all 88. The executed T1 source is byte-identical between the initial and EIA phases (SHA-256 `df1950ef3930...d56d522`). Reporting-period identities are a separate representation layer: the initial replay uses canonical month/quarter identities for relevant families, while EIA uses day-level data-ending dates. These reporting-period strings never enter T1’s release-date comparator and are adjudicated by the frozen family scorer.

A.7 COMPLETED IDENTIFICATION LADDER

Table 4 separates four distinct questions that are all executed on the source-native replay. The controls were added prospectively in sequence: outcome thresholds, endpoint units, integration surfaces, model identities, and checkpoint policies were frozen before the corresponding fresh outcomes.

G₀: same-surface no-operation placebo. G₀ executes exactly one assigned helper call through the same harness surface as T but is information-empty: the frozen function is `def skill(package, context): return {}`. It reads neither package nor context. The first DeepSeek 35-endpoint two-repeat tranche gives G₀-N 90% CI $[-14.3, 2.9]$ points; an independently frozen four-repeat A2 study gives -2.1 points with 90% CI $[-9.3, 5.0]$ over the full endpoint portfolio, and Kimi gives $+2.2$ points with CI $[-4.3, 8.7]$. These tranches are never pooled. A post-review exact-stratum audit shows why full-track equivalence should not be transported mechanically: A2 C3/C4 grounding each have G₀-N $=+5$ points with 90% intervals reaching $+15$, while the eight EIA held-out endpoints have -18.8 points with a wide interval. We therefore treat G₀-N as a surface-perturbation diagnostic. In

Question	Contrast	Status	Identified quantity
Repair original agent?	T-N	Executed	Incremental value of inserting the targeted procedure over the original agent.
Beat target-blind help?	T-G	Executed	Targeted-operation value relative to the frozen complexity-matched helper.
Does operation output matter on the same surface?	T-G ₀	Executed	Same-surface operation-output contrast; G ₀ -N is reported separately as the placebo/surface perturbation.
Does integration surface matter here?	T-R _{surf}	Executed	Effect of placing the exact same precomputed operation output on the T surface versus ordinary context under the forced one-answer harness.

Table 4: Completed identification ladder. G₀ isolates operation output under the same helper surface without assuming stratum-level neutrality; R_{surf} isolates the integration surface while holding operation output exactly fixed.

the fresh A1 tranche, C3 and C4 grounding have G₀=N exactly and T exceeds both by +30 and +20 points; EIA has T-N=+37.5 and T-G₀=+62.5, but G₀-N=-25, so the EIA net-repair claim is carried by T-N rather than a neutrality assumption.

R_{surf}: exact-output context-materialization control. R_{surf} runs the exact same targeted operation source as T on the same candidate evidence. Rather than exposing the output through `reusable_helper_output`, the harness adds the identical object as `retrieval_operation_output` inside the evidence package and exposes no callable helper. A zero-call parity audit verifies for every tested endpoint that (i) removing this one field recovers the original evidence package exactly, (ii) operation-output content hashes are identical between T and R_{surf}, and (iii) no candidate evidence is removed or reordered. On all 18 pre-frozen DeepSeek load-bearing endpoints, two fresh repeats per arm give T-R_{surf} = 0.0 points with bootstrap 90% CI [-5.6, 5.6] and post-hoc 95% sensitivity CI [-8.3, 8.3]. A paired-endpoint TOST at the prospectively frozen ± 10 -point margin gives $p = 0.0121$ on each one-sided test; paired-t 95% CI is [-8.5, 8.5]. A deterministic parametric sensitivity calculation using the observed endpoint SD (17.15 points) gives about 53% probability of declaring equivalence when the true mean is exactly zero, so we do not characterize the design as high-powered. C3 grounding is T=R_{surf}=60% with five endpoint ties; C4 grounding is 70%/70% with one win, three ties, and one loss; EIA is ceiling at 100%/100%. The four strictly non-ceiling fresh endpoints have mean residual 0 but a wide 90% interval [-25, 25]. A secondary Kimi alignment check also ties on all runs. We therefore describe the observed DeepSeek portfolio as consistent with equivalence at the ± 10 -point *portfolio-average* margin under this forced one-answer harness, with low sensitivity and a ceiling-heavy composition. Fine-grained non-ceiling effects, adaptive callable use, and independently designed temporal retrievers remain outside the claim.

A.8 R_{surf} ENDPOINT RESIDUALS AND FUTURE RESOLUTION PLANNING

The ± 10 -point threshold was frozen before R_{surf} outcomes as a materiality threshold and was not derived from the observed T-R_{surf} variance. The pre-outcome contract does not record a variance-based power rationale. Post review, we note only as contextual scale information that targeted contrasts already visible before the R_{surf} run were on a 20–60 point scale; this does not retroactively turn that scale argument into a preregistered justification or make the completed $n = 18$ design high-powered. Using the observed pooled endpoint SD of 17.15 points, a deterministic post-review planning simulation gives about 53% probability of passing the same TOST when the true mean is zero at $n = 18$, and about 81% at $n = 27$. The latter is a future-design estimate under the same variance, not retroactive evidence. Family-specific variance estimates are too unstable to justify smaller targets: C3 grounding and EIA have zero observed residual variance because all endpoint means tie (EIA also ceilings), whereas the five C4 grounding endpoints have SD 35.36 points; under that observed variance, even $n = 100$ gives only about 75% simulated sensitivity and roughly $n = 110$ is needed to reach 81%. Future non-ceiling work should therefore pre-register an independent endpoint count and stopping rule rather than expand until an equivalence test passes.

The two non-ties do not point in the same direction. On endpoint 15ffceb0, only T repeat 0 selects the registered successful span pair; T repeat 1 and both R_{surf} repeats select the same failing

Endpoint suffix	Stratum	T (%)	R _{surf} (%)	Δ (pp)
128dc1bd	C3 Grd	100	100	0
15ffceb0	C4 Grd	50	0	+50
187e543f	C4 Grd	0	50	-50
1dfccd3a	C3 Grd	0	0	0
5c43ed3f	C4 Grd	100	100	0
65c03428	C3 Grd	100	100	0
70e1e3d1	C3 Grd	0	0	0
71b83aed	C4 Grd	100	100	0
95037898	C4 Grd	100	100	0
cc293f19	C3 Grd	100	100	0
229defbd	EIA Cut	100	100	0
4054808c	EIA Cut	100	100	0
62014739	EIA Cut	100	100	0
9bf3bfa4	EIA Cut	100	100	0
a4f03bdc	EIA Cut	100	100	0
b246cd01	EIA Cut	100	100	0
c0799777	EIA Cut	100	100	0
d6a5cf91	EIA Cut	100	100	0

Table 5: All 18 pre-frozen DeepSeek endpoint-mean R_{surf} residuals. Endpoint suffixes are shortened display identifiers; full IDs and repeat-level rows are in the supplement.

pair. On 187e543f, only R_{surf} repeat 0 succeeds; both T repeats and R_{surf} repeat 1 select the same failing pair. The targeted operation-output object is content-identical across T and R_{surf} for both endpoints. The opposite-signed, single-repeat deviations therefore do not isolate a persistent one-direction surface mechanism.

A.9 PROSPECTIVE EXTENSION TO ADAPTIVE MULTI-TURN SURFACES

The current R_{surf} result is intentionally limited to a forced one-answer harness. A multi-turn extension would be a new experiment, not an inference from the present rows. A valid design should (i) use independent episodes, not turns, as scientific units; (ii) freeze the candidate-evidence stream and operation implementation before outcomes; (iii) separate an *availability* estimand (callable interface present versus absent) from a *placement* estimand in which the exact same operation output is synchronized across callable and context-materialized arms; and (iv) pre-register how endogenous invocation opportunities are aligned, rather than selecting turns after observing successful calls. A same-surface no-op arm can diagnose interface perturbation at each registered opportunity. This would test whether adaptive callable state matters when invocation and context evolve across turns; the present study makes no such claim.

A benign mechanism-agnostic generic helper would likewise be a new treatment. It could be useful for triangulating typical T-versus-generic practice, but it is not retroactively inserted into the current identification ladder: G remains a stress comparator and G₀ a same-surface no-operation placebo. Any future benign-generic study must freeze its behavior, endpoint units, and decision rule before new outcomes.

A.10 FOLLOW-UP EXTENSION: BENIGN ORGANIZATION, CROSS-DOMAIN REPLAY, AND ACTUAL TOOL USE

After the primary audit and its claim boundary were frozen, we ran a separate follow-up extension. These outcomes were *not* used to choose the primary claim and are not retroactively described as preregistered evidence for it. Their role is narrower: test whether a more realistic mechanism-agnostic helper absorbs targeted credit, obtain fresh non-ceiling and cross-domain endpoints, and check whether real multi-turn function availability creates a residual that the forced one-answer R_{surf} experiment could not measure. All extension calls use the same provider route and exact DeepSeek model identity as the primary fresh controls; each completed call or turn is checkpointed before the next unit. Repeats are averaged within endpoint and never increase the scientific sample size.

A benign generic evidence organizer. We introduce a new arm B that is deliberately useful but mechanism-agnostic. B copies and organizes the complete supplied evidence in its original order and adds a generic verification reminder. It is forbidden from filtering by release date, inferring reporting

periods, selecting mechanism-specific grounding spans, or reordering candidates. Thus B differs from the historical mechanism-misaligned stress helper G and from the information-empty placebo G_0 . An initial four-endpoint historical pilot was all ties and therefore not scaled. We then froze genuinely new EIA releases and a new BLS CPI substrate before their respective model outcomes. A first-party EIA archive page dated 2026-06-24 contains an internally inconsistent release-date header; that source anomaly was excluded before model outcomes rather than silently corrected. The clean EIA targets are the July 22, July 29, August 5, and August 12 releases, with the next release used only as a post-cutoff distractor. The BLS set uses four CPI release targets from April through July 2026 and first-party archive-header fields matching the information shown to the model.

Fresh substrate	N	B	T	T-N	B-N	T-B
EIA WPSR ($n = 4$)	.375	1.000	1.000	+.625	+.625	.000
BLS CPI ($n = 4$)	.375	.750	1.000	+.625	+.375	+.250

Table 6: Follow-up extension using endpoint-mean family success after two execution repeats per endpoint. Repeats measure robustness and are averaged within the four independent endpoints. B is the mechanism-agnostic evidence organizer.

The EIA result is especially diagnostic: T still repairs the original agent by 62.5 points on average, but B produces the same endpoint means as T on all four endpoints. On this substrate the targeted operation therefore receives no residual credit over generic evidence organization. BLS is mixed rather than uniformly absorbed: generic organization explains part of the net repair, while one target has stable $T > B = N$ across both repeats. The two fresh sets therefore sharpen, rather than overturn, the paper’s attribution message: T-N can establish net repair without establishing that the registered temporal operation uniquely deserves that credit. Because each set has only four endpoints, these follow-ups are robustness and boundary evidence rather than high-resolution effect estimates.

Actual multi-turn function availability. We separately verified that the provider supports a genuine function-call round trip: the model emits a function call, receives a `function_call_output`, and then produces a final answer. On the four clean EIA endpoints, each experimental unit begins from one shared first-turn planning response and then forks into BASE (no helper), CONTEXT (the exact frozen T1 output materialized as ordinary context), or TOOL (a real `temporal_filter` function is available and invocation is endogenous). All three arms succeed on 4/4 endpoints in this exploratory fork. TOOL is actually invoked on 2/4 endpoints, yet TOOL-CONTEXT is zero on every endpoint. This is a null boundary rather than evidence for multi-turn equivalence: the planning turn itself changes the baseline, the design is small, and endogenous invocation is part of the TOOL treatment, so TOOL-CONTEXT is not a pure placement estimand.

A planning-baseline hypothesis fails prospectively. The unexpectedly strong helper-free BASE arm suggested a possible spin-off hypothesis: perhaps an extra structured planning/decomposition turn absorbs apparent temporal-skill gains. On the EIA endpoints, two planning-only executions give endpoint-mean success .875 versus .375 for the matched single-turn N runs; on the BLS endpoints the corresponding exploratory values are 1.000 versus .375. These observations motivated, but do not confirm, the hypothesis: the EIA result generated the idea, and the BLS planning arm was added after the original BLS outcomes were visible.

We therefore froze a third-substrate falsifier *before any Federal Reserve model outcome*. Four FOMC statement targets (December 10, 2025; January 28, March 18, and April 29, 2026) were selected from one consecutive first-party statement chain, each with three prior statements and one next-statement distractor. The comparison was fixed as single-turn N versus helper-free two-turn planning before either arm ran. Single-turn N succeeds on 1/4 endpoints and planning-only also succeeds on 1/4. Pairwise, planning has one win, N has one win, and two endpoints tie, for mean difference zero. Inspection shows no stable cutoff mechanism: on one endpoint planning removes the future distractor, while on another it introduces one that the single-turn answer had excluded. We therefore *stop* the proposed planning/deliberation paper direction at this point. The positive EIA/BLS observations remain visible exploratory evidence, but they do not survive the prospectively frozen independent-substrate falsifier. Reopening that scientific object would require a new prospectively frozen substrate showing a stable planning-over-single-turn residual under the same-information contract.

Checkpointed execution and recovery. Every fresh G_0 and R model call is persisted before the next unit begins: an atomic raw JSON receipt, one fsynced CSV row, one fsynced JSONL row, and an atomic checkpoint summary keyed by endpoint $\times$ repeat $\times$ arm. Resume logic executes missing units only and never reruns a successfully checkpointed unit. Each experiment first runs a small runtime/protocol pilot whose promotion rule ignores scientific outcomes. This design was exercised in practice: transient control-channel failures occurred during the Kimi run while the remote process continued, and the already completed rows remained recoverable from CSV/raw checkpoints without replaying prior calls.

Every retained row binds endpoint, condition, repeat, skill hash, prompt hash, raw response hash, and provider response-ID hash. Kimi repeat 4 is excluded by hash provenance with no replacement. The exact TimeSage-EV replication remains separate from all source-native evidence.